HyperTransport MegaCore Function

User Guide

Contents

Release Information

Table 1–1 provides information about this release of the HyperTransport MegaCore® function.

Table 1–1. HyperTransport MegaCore Function Release Information

Item	Description
Version	9.0
Release Date	March 2009
Ordering Code	IP-HT
Product ID(s)	0098
Vendor ID(s)	6AF7

Altera® verifies that the current version of the Quartus® II software compiles the previous version of each MegaCore function. Any exceptions to this verification are reported in the *MegaCore IP Library Release Notes and Errata*. Altera does not verify compilation with MegaCore function versions older than one release.

Device Family Support

MegaCore functions provide either full or preliminary support for target Altera device families:

■ *Full support* means the MegaCore function meets all functional and timing requirements for the device family and may be used in production designs.

■ *Preliminary support* means the MegaCore function meets all functional requirements, but may still be undergoing timing analysis for the device family; it may be used in production designs with caution.

Table 1–2 shows the level of support offered by the HyperTransport MegaCore function for each of the Altera device families.

Table 1–2. Device Family Support

Device Family	Support
HardCopy® Stratix®	Full
Stratix	Full
Stratix II	Full
Stratix II GX	Preliminary
Stratix GX	Full
Other device families	No support

Introduction

The HyperTransport MegaCore function implements high-speed packet transfers between physical (PHY) and link-layer devices, and is fully compliant with the HyperTransport I/O Link Specification, Revision 1.03. This MegaCore function allows designers to interface to a wide range of HyperTransport™ technology (HT) enabled devices quickly and easily, including network processors, coprocessors, video chipsets, and ASICs.

Features

The HyperTransport MegaCore function has the following features:

- 8-bit fully integrated HT end-chain interface
- Packet-based protocol
- Dual unidirectional point-to-point links
- Up to 16 Gigabits per second (Gbps) throughput (8 Gbps in each direction)
    - 200, 300, and 400 MHz DDR links in Stratix and Stratix GX devices
    - 200, 300, 400, and 500 MHz DDR links in Stratix II and Stratix II GX devices
- Low-swing differential signaling with 100-Ω differential impedance
- Hardware verified with HyperTransport interfaces on multiple industry standard processor and bridge devices
- Fully parameterized MegaCore function allows flexible, easy configuration
- Fully optimized for the Altera Stratix II, Stratix, Stratix GX, and Stratix II GX device families
- Application-side interface uses the Altera Atlantic™ interface standard
- Manages HT flow control, optimizing performance and ease of use
- Independent buffering for each HT virtual channel
    - Automatic handling of HT ordering rules
    - Stalling of one virtual channel does not delay other virtual channels (subject to ordering rules)
    - Flexible parameterized buffer sizes, allowing customization depending on system requirements
    - User interface has independent interfaces for the HT virtual channels, allowing independent user logic design
- Cyclic redundancy code (CRC) generation and checking to preserve data integrity
- Integrated detection and response to common HT error conditions
    - CRC errors
    - End-chain errors
- Fully integrated HT configuration space includes all required configuration space registers and HT capabilities list registers

■ 32-bit and 64-bit support across all base address registers (BARs)

■ Automatically handles all CSR space accesses

■ Verilog HDL and VHDL simulation support

OpenCore Plus Evaluation

With the Altera free OpenCore Plus evaluation feature, you can perform the following actions:

■ Simulate the behavior of a megafunction (Altera MegaCore function or AMPP™ megafunction) within your system

■ Verify the functionality of your design, as well as quickly and easily evaluate its size and speed

■ Generate time-limited device programming files for designs that include MegaCore functions

■ Program a device and verify your design in hardware

You only need to purchase a license for the MegaCore function when you are completely satisfied with its functionality and performance, and want to take your design to production.

For more information about OpenCore Plus hardware evaluation using the HyperTransport MegaCore function, refer to "OpenCore Plus Time-Out Behavior" on page 3–40 and *AN 320: OpenCore Plus Evaluation of Megafunctions*.

Performance

The HyperTransport MegaCore function uses 20 differential I/O pin pairs and 2 single-ended I/O pins, requiring 42 pins total. Table 1–3 through Table 1–5 show typical performance and adaptive look-up table (ALUT) or logic element (LE) usage for the HyperTransport MegaCore function in Stratix II GX, Stratix II, Stratix, and Stratix GX devices respectively, using the Quartus® II software version 7.1.

Table 1–3 shows the maximum supported data rates in megabits per second (Mbps) by device family and speed grade.

Table 1–3. Maximum Supported HyperTransport Data Rates *(Note 1)*

Device Family	Speed Grade					
	-3	**-4**	**-5**	**-6**	**-7**	**-8**
Stratix II GX devices	1000 Mbps	1000 Mbps	800 Mbps	N/A *(2)*	N/A *(2)*	N/A *(2)*
Stratix II devices	1000 Mbps	1000 Mbps	800 Mbps	N/A *(2)*	N/A *(2)*	N/A *(2)*
Stratix devices (Flip-Chip packages)	N/A *(2)*	N/A *(2)*	800 Mbps	800 Mbps	600 Mbps	400 Mbps
Stratix devices (Wire Bond packages)	N/A *(2)*	N/A *(2)*	N/A *(2)*	600 Mbps	400 Mbps	400 Mbps
Stratix GX devices	N/A *(2)*	N/A *(2)*	800 Mbps	800 Mbps	600 Mbps	N/A *(2)*

Notes to Table 1–3:

(1) Rates are per interface bit. Multiply by eight to calculate the uni-directional data rate of an 8-bit interface.

(2) Devices of this speed grade are not offered in this device family.

Table 1–4 shows performance and device utilization for the HyperTransport MegaCore function in Stratix II and Stratix II GX devices.

Table 1–4. HyperTransport MegaCore Function Performance in Stratix II and Stratix II GX Devices

| Parameters | | | | Combinational ALUTs *(2)* | Logic Registers | Memory | | HT Link f_{MAX}(MHz) *(3)* | User Interface f_{MAX} (MHz) *(3)* |
Rx Posted Buffers	Rx Non-Posted Buffers	Rx Response Buffers	Clocking Option *(1)*			M4K	M512		
8	4	4	Shared Rx/Tx/Ref	3,500	5,200	12	0	500	125 *(4)*
8	4	4	Shared Ref/Tx	3,500	5,200	14	0	500	125 *(4)*
8	4	4	Shared Rx/Tx	3,600	5,400	16	0	500	> 150
8	8	8	Shared Rx/Tx	4,000	6,000	16	0	500	> 150
16	8	8	Shared Rx/Tx/Ref	4,100	6,200	12	0	500	125 *(4)*
16	8	8	Shared Ref/Tx	4,100	6,200	14	0	500	125 *(4)*
16	8	8	Shared Rx/Tx	4,200	6,400	16	0	500	> 150

Notes to Table 1–4:

(1) Refer to "Clocking Options" on page 3–7 for more information about these options.

(2) Other parameters (BAR configurations, etc.) vary the ALUT and Logic Register utilization numbers by approximately +/- 200.

(3) Figures for -3 speed grade devices only.

(4) When using the **Shared Rx/Tx/Ref** and **Shared Ref/Tx** options, the user interface frequency is limited to exactly the HT frequency divided by four.

Table 1–5 shows performance and device utilization for the HyperTransport MegaCore function in Stratix and Stratix GX devices.

Table 1–5. HyperTransport MegaCore Function Performance in Stratix and Stratix GX Devices

| Parameters | | | | Utilization | | HT Link f_{MAX}(MHz) | | User Interface f_{MAX} (MHz) | |
| Rx Posted Buffers | Rx Non-Posted Buffers | Rx Response Buffers | Clocking Option *(1)* | LEs *(2)* | M4K Blocks | Speed Grade | | | |
						-5	-6	-5	-6
8	4	4	Shared Rx/Tx/Ref	7,500	12	400	400	100 *(3)*	100 *(3)*
8	4	4	Shared Ref/Tx	7,600	14	400	400	100 *(3)*	100 *(3)*
8	4	4	Shared Rx/Tx	7,900	16	400	400	> 125	> 100
8	8	8	Shared Rx/Tx	8,900	16	400	400	> 125	> 100
16	8	8	Shared Rx/Tx/Ref	9,400	12	400	400	100 *(3)*	100 *(3)*
16	8	8	Shared Ref/Tx	9,500	14	400	400	100 *(3)*	100 *(3)*
16	8	8	Shared Rx/Tx	9,700	16	400	400	> 125	> 100

Notes to Table 1–5:

(1) Refer to "Clocking Options" on page 3–7 for more information about these options.

(2) Other parameters (BAR configurations etc.) vary the LE utilization by approximately +/- 200 LEs.

(3) When using the **Shared Rx/Tx/Ref** and **Shared Ref/Tx** options, the user interface frequency is limited to exactly the HT frequency divided by four.

Design Flow

To evaluate the HyperTransport MegaCore function using the OpenCore Plus feature, include these steps in your design flow:

1. Obtain and install the HyperTransport MegaCore function.

The HyperTransport MegaCore function is part of the MegaCore IP Library, which is distributed with the Quartus II software and downloadable from the Altera website, www.altera.com.

For system requirements and installation instructions, refer to *Quartus II Installation & Licensing for Windows and Linux Workstations* on the Altera website at www.altera.com/literature/lit-qts.jsp.

Figure 2–1 shows the directory structure after you install the HyperTransport MegaCore function, where *<path>* is the installation directory. The default installation directory on Windows is **C:\altera\<version number>**; on Linux it is **/opt/altera<version number>**.

Figure 2–1. Directory Structure
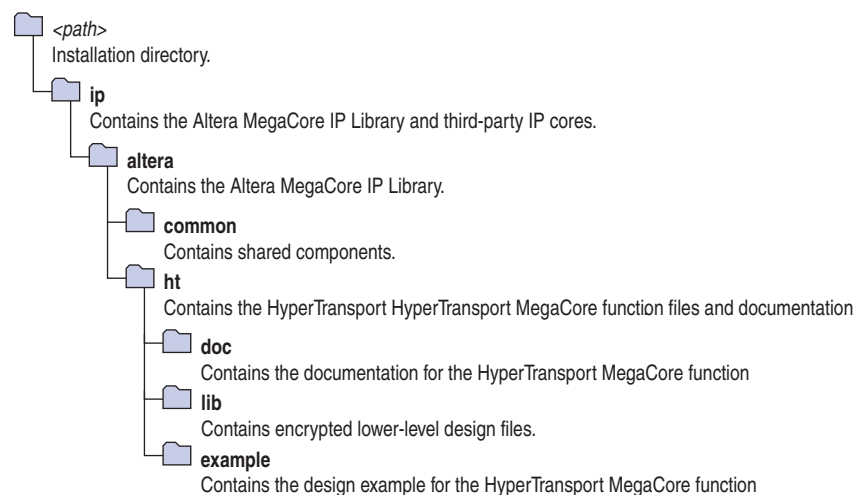
2. Create a custom variation of the HyperTransport MegaCore function.

3. Implement the rest of your design using the design entry method of your choice.

4. Use the IP functional simulation model to verify the operation of your design.

For more information about IP functional simulation models, refer to the *Simulating Altera IP in Third-Party Simulation Tools* chapter in volume 3 of the *Quartus II Handbook*.

5. Use the Quartus II software to compile your design.

☞ You can also generate an OpenCore Plus time-limited programming file, which you can use to verify the operation of your design in hardware.

6. Purchase a license for the HyperTransport MegaCore function.

After you have purchased a license for the HyperTransport MegaCore function, follow these additional steps:

1. Set up licensing.

2. Generate a programming file for the Altera device(s) on your board.

3. Program the Altera device(s) with the completed design.

MegaCore Function Walkthrough

This walkthrough explains how to create a custom variation using the Altera HyperTransport IP Toolbench and the Quartus II software, and simulate the function using an IP functional simulation model and the ModelSim software. When you are finished generating your custom variation of the function, you can incorporate it into your overall project.

☞ IP Toolbench allows you to select only legal combinations of parameters, and warns you of any invalid configurations.

In this walkthrough, you follow these steps:

■ Create a New Quartus II Project

■ Launch the MegaWizard Plug-in Manager

■ Step 1: Parameterize

■ Step 2: Set Up Simulation

■ Step 3: Generate

■ Simulate the Design

☞ To generate a wrapper file and IP functional simulation model using default values, omit the procedure described in "Step 1: Parameterize" on page 2–5.

Create a New Quartus II Project

Create a new Quartus II project with the **New Project Wizard**, which specifies the working directory for the project, assigns the project name, and designates the name of the top-level design entity.

To create a new project, perform the following steps:

1. On the Windows Start menu, select **Programs > Altera > Quartus II** *<version>* to start the Quartus II software. Alternatively, you can use the Quartus II Web Edition software.

2. In the Quartus II window, on the File menu, click **New Project Wizard**. If you did not turn it off previously, the **New Project Wizard: Introduction** page appears.

3. On the **New Project Wizard Introduction** page, click **Next**.

4. On the **New Project Wizard: Directory, Name, Top-Level Entity** page, enter the following information:

   a. Specify the working directory for your project. For example, this walkthrough uses the **C:\altera\projects\ht_project** directory.

   b. Specify the name of the project. This walkthrough uses `ht_example` for the project name.

   ☞ The Quartus II software automatically specifies a top-level design entity that has the same name as the project. This walkthrough assumes that the names are the same.

5. Click **Next** to close this page and display the **New Project Wizard: Add Files** page.

   ☞ When you specify a directory that does not already exist, a message asks if the specified directory should be created. Click **Yes** to create the directory.

6. Click **Next** to close this page and display the **New Project Wizard: Family and Device Settings** page.

7. On the **New Project Wizard: Family & Device Settings** page, perform the following steps:

   a. in the **Family** list, select the target device family.

   b. Under **Target device**, turn on **Specific device selected in 'Available devices' list**.

   c. In the **Available devices** list, select a device.

8. The remaining pages in the **New Project Wizard** are optional. Click **Finish** to complete the Quartus II project.

You have finished creating your new Quartus II project.

Launch the MegaWizard Plug-in Manager

To launch the MegaWizard™ Plug-in Manager in the Quartus II software, perform the following steps:

1. On the **Tools** menu, click **MegaWizard Plug-In Manager**. The MegaWizard Plug-In Manager displays, as shown in Figure 2–2.

   ☞ Refer to Quartus II Help for more information about how to use the MegaWizard Plug-In Manager.

Figure 2–2. MegaWizard Opening Screen
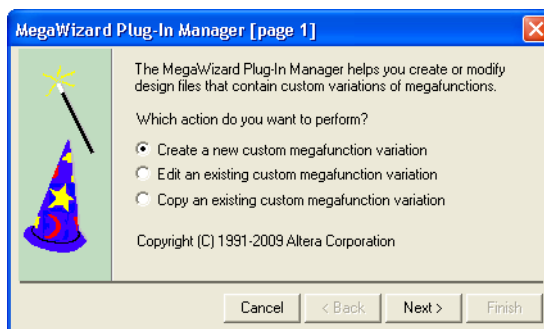
2. Choose **Create a new custom megafunction variation** and click **Next**.

3. Under **Interfaces** in the **HyperTransport** folder, click the **HT v9.0** component.

4. Choose the device family you want to use for this MegaCore function variation, for example, **Stratix II GX**. Your selection should match the device family you selected in step 7 on page 2–3 when creating the project.

5. Select the output file type for your design; the wizard supports VHDL and Verilog HDL.

6. The MegaWizard Plug-in Manager shows the project path that you specified in the New Project Wizard. Append a variation name for the MegaCore function output files *<project path>\<variation name>*. For this walkthrough, to create a project that includes only a single HyperTransport MegaCore function with no additional logic, define *<variation name>* to be `ht_example` to match the project name. Figure 2–3 shows the wizard after you have made these settings.

Figure 2–3. Select the MegaCore Function
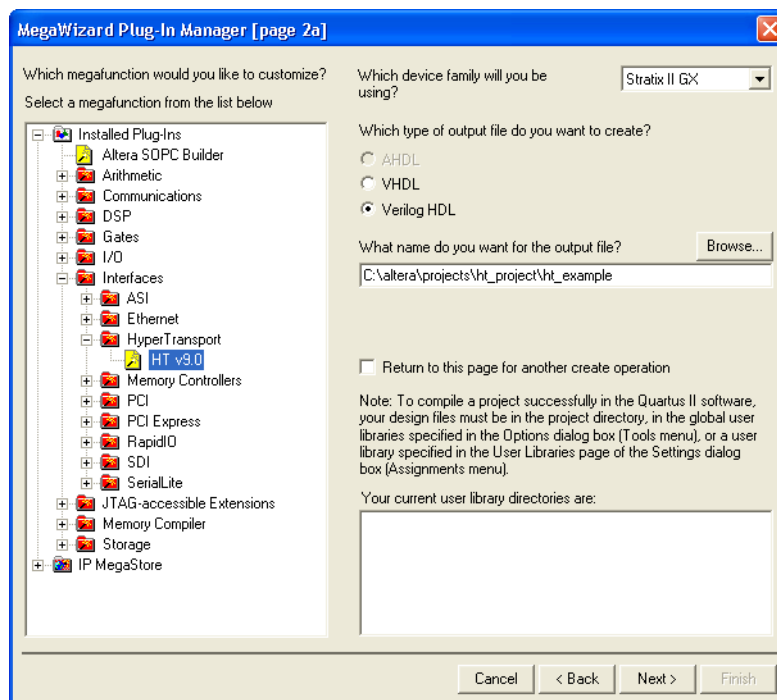
7. Click **Next** to launch the IP Toolbench for the HyperTransport MegaCore function.

Figure 2–4. IP Toolbench

Step 1: Parameterize

To parameterize your MegaCore function, follow these steps:

1. In the IP Toolbench, click **Step 1: Parameterize**, as shown in Figure 2–4. The IP Toolbench wizard opens to the **Device Family & Read-Only Registers** tab.

2. Set the target Altera device family and the values for the read-only HyperTransport configuration registers on the **Device Family & Read-Only Registers** tab. For this walkthrough, use the default settings, which are shown in Figure 2–5. For more information about these parameters, refer to Table A–1 on page A–1.

Figure 2–5. Parameterize—Device Family and Read-Only Registers

3. Click **Next**. The **Base Address Registers** tab appears.

4. On the **Base Address Registers** tab, configure the HyperTransport BARs that define the address ranges of memory read and write request packets that your application claims from the HyperTransport interface. For this walkthrough, use the default settings, which are shown in Figure 2–6. For more information about the parameters modified by these settings, refer to Table A–2 on page A–2.

Figure 2–6. Parameterize—Base Address Registers Tab
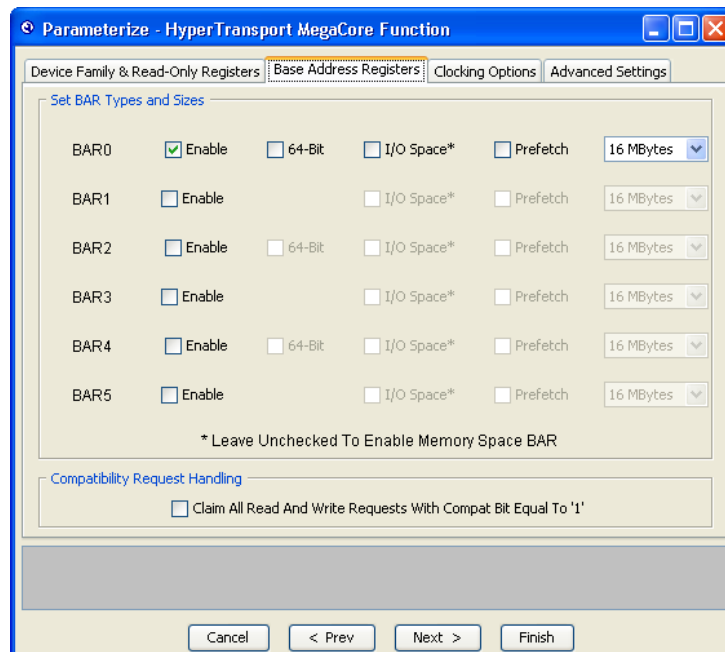
5. Click **Next**. The **Clocking Options** tab displays, as shown in Figure 2–7.

6. You set the clocking options for your application on the **Clocking Options** tab. For more information about the available options, refer to "Clocking Options" on page 3–7 and Table A–3 on page A–3. For this walkthrough, use the default settings, which are shown in Figure 2–7.

☞ HyperTransport link clock frequencies of 500 MHz are only supported in Stratix II and Stratix II GX devices.

Figure 2–7. Parameterize—Clocking Options Tab
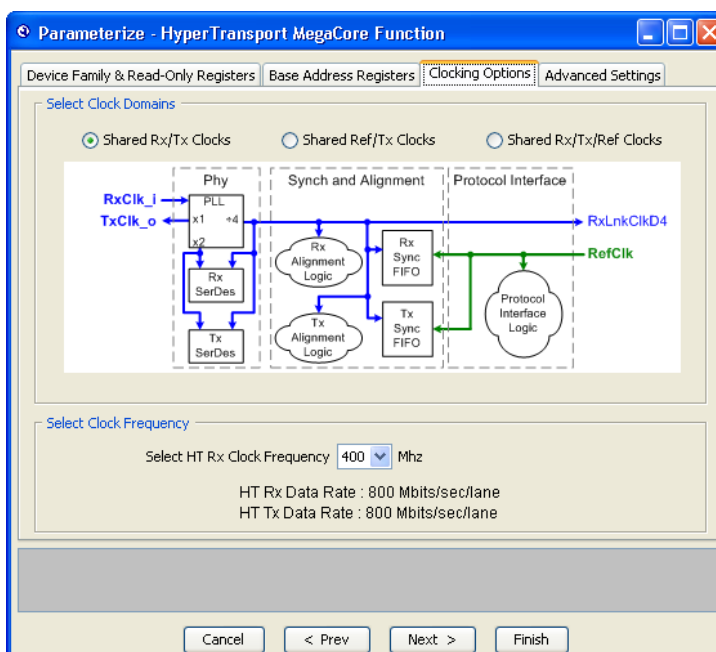
7. Click **Next**. The **Advanced Settings** tab displays.

8. You set the receiver virtual channel buffer sizes and the maximum allowed delay from the deassertion of `TxRDav_o` to the assertion of `TxRSop_i` on the **Advanced Settings** tab. For more information about these parameters, refer to Table A–4 on page A–3. For this walkthrough, use the default settings, which are shown in Figure 2–8.

Figure 2–8. Parameterize—Advanced Settings Tab
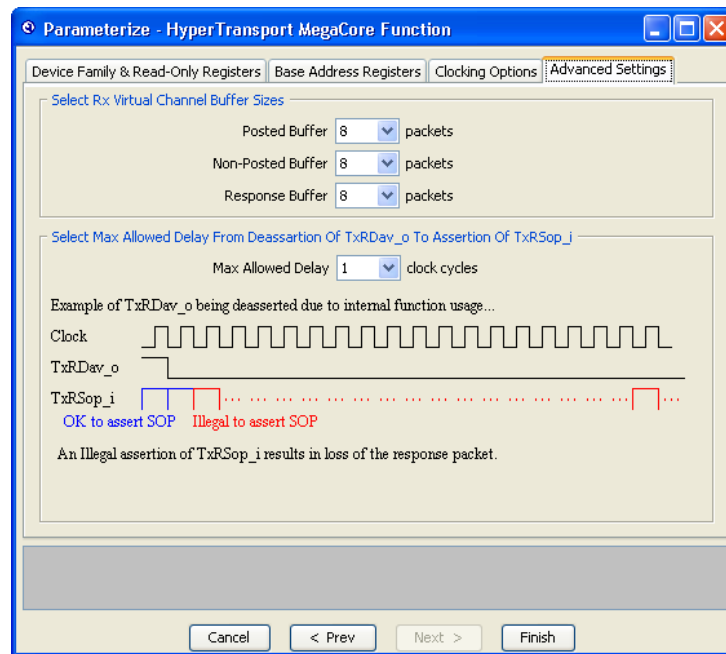
9. Click **Finish**. The **Parameterize—HyperTransport MegaCore Function Parameterize** panel closes and you are returned to the IP Toolbench interface.

Step 2: Set Up Simulation

An IP functional simulation model is a cycle-accurate VHDL or Verilog HDL model file produced by the Quartus II software. The simulation model allows for fast functional simulation of IP using industry-standard VHDL and Verilog HDL simulators.

☞ You may only use these simulation model output files for simulation purposes and expressly not for synthesis or any other purposes. Using these models for synthesis creates a nonfunctional design.

To generate an IP functional simulation model for your MegaCore function, follow these steps:

1. In the IP Toolbench, click **Step 2: Set Up Simulation**, as shown in Figure 2–9.

Figure 2–9. Set Up Simulation

2. Turn on **Generate Simulation Model**, as shown in Figure 2–10.

Figure 2–10. Generate Simulation Model

3. Select the language in the **Language** list. In this case, **Verilog HDL** was chosen.

If you are synthesizing your design with a third-party EDA synthesis tool, you can generate a netlist for the synthesis tool to estimate timing and resource usage for this megafunction.

1. To generate a netlist, turn on **Generate netlist**.

2. Click **OK**.

Step 3: Generate

To generate your MegaCore function, follow these steps:

1. In the IP Toolbench, click **Step 3: Generate** as shown in Figure 2–11.

Figure 2–11. Generation
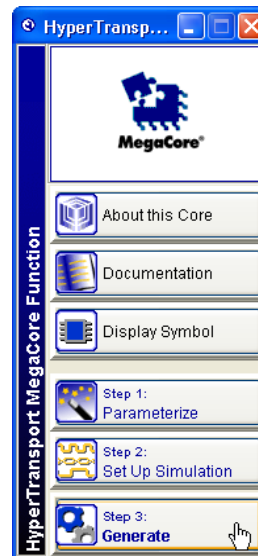
The generation report lists the design files that the IP Toolbench creates, as shown in Figure 2–12.

Figure 2–12. Generation Report—HyperTransport MegaCore Function

2. After the MegaCore function is generated, according to the message and progress at the bottom of the generation report window, click **Exit**.

3. If you are prompted to add the Quartus II IP File (**.qip**) to the project, click **Yes**.

☞ If you previously turned on **Automatically add Quartus II IP Files to all projects**, the **.qip** file is generated automatically.

You have generated an instance of the HyperTransport MegaCore function.

Table 2–1 describes the IP Toolbench-generated files, which are listed in the file *<variation name>*.**html** in your project directory.

Table 2–1. IP Toolbench Files *(Note 1)*

File Name (2)	Description
<variation name>.**vhd** or .**v**	A MegaCore function variation file, which defines a VHDL or Verilog HDL top-level description of the custom MegaCore function. Instantiate the entity defined by this file inside your design. Include this file when compiling your design in the Quartus II software.
*<variation name>*_**bb.v**	Verilog HDL black-box file for the MegaCore function variation. Use this file when using a third-party EDA tool to synthesize your design. This file is only produced when the Verilog HDL language is selected.
<variation name>.**bsf**	Quartus II symbol file for the MegaCore function variation. You can use this file in the Quartus II block diagram editor.
<variation name>.**cmp**	A VHDL component declaration file for the MegaCore function variation. Add the contents of this file to any VHDL architecture that instantiates the MegaCore function. This file is only produced when the VHDL language is selected.
<variation name>.**vo** or *<variation name>*.**vho**	VHDL or Verilog HDL IP functional simulation model.
<variation name>.**qip**	Contains Quartus II project information for your MegaCore function variation.
<variation name>.**html**	The MegaCore function report file.

Notes to Table 2–1:

(1) These files are variation dependent; some may be absent or their names may change.

(2) *<variation name>* is the variation name selected by the user in the MegaWizard Plug-In Manager.

☞ The **.qip** file is generated by the MegaWizard interface and contains information about your generated IP core. You are prompted to add this **.qip** file to the current Quartus II project at the time of file generation. In most cases, the **.qip** file contains all of the necessary assignments and information required to process the core or system in the Quartus II compiler. Generally, a single **.qip** file is generated for each MegaCore function.

You can now integrate your custom megafunction in your design and compile the design.

Simulate the Design

To simulate your design, you use the IP functional simulation models generated by the IP Toolbench. The IP Functional Simulation model is the **.vo** or **.vho** file generated by the IP Toolbench, as specified in "Step 2: Set Up Simulation" on page 2–9. Add this file in your simulation environment to perform functional simulation of your custom variation of the MegaCore function.

The HyperTransport MegaCore function vector-based testbench is an example you can use to help set up your own simulation environment. You should not attempt to edit these files. For information about how to perform a simulation using this vector-based testbench, see "Example Simulation and Compilation" on page 2–16.

For more information about IP functional simulation models, refer to the *Simulating Altera IP in Third-Party Simulation Tools* chapter in volume 3 of the *Quartus II Handbook*. You can use any Altera-supported third-party simulator to simulate your design and testbench.

Compile the Design

You can use the Quartus II software to compile your design. Refer to Quartus II Help for instructions on compiling your design.

The instructions in this section assume that you named your wrapper file **ht_example.v** using the MegaWizard Plug-In Manager. If you chose a different name, substitute that name when following the instructions.

To compile your design in the Quartus II software, perform the following steps:

1. If you are using the Quartus II software to synthesize your design, skip to step 2. If you are using a third-party synthesis tool to synthesize your design, perform the following steps:

   a. Set a black box attribute for **ht_example.v** before you synthesize the design. Refer to the Quartus II Help for your specific synthesis tool for instructions on setting black-box attributes.

   b. Run the synthesis tool to produce an EDIF Netlist File (**.edf**) or Verilog Quartus Mapping file (**.vqm**) for input to the Quartus II software.

   c. Add the **.edf** or .vqm file to your Quartus II project.

2. On the **Processing** menu, point to **Start** and click **Start Analysis & Elaboration** to elaborate the design.

3. On the Assignments menu, click **Assignment Editor**.

4. If the pin names are not displayed, on the View menu, click **Show All Known Pin Names**.

5. Set the **I/O Standard** to **HyperTransport** for the I/O pins that are connected to the HyperTransport wrapper ports `TxCAD_o[7:0]`, `TxCTL_o`, `TxClk_o`, `RxCAD_i[7:0]`, `RxCTL_i`, and `RxClk_i`, by performing the following steps:

   a. In the row for the pin, double-click in the **Assignment Name** column.

   b. In the Assignment Name list, click **I/O Standard**.

   c. In the row for the pin, double-click in the **Value** column.

   d. In the Value list, click **HyperTransport**.

6. Set the **I/O Standard** to **2.5 V** for the I/O pins connected to the HyperTransport wrapper ports `PwrOk` and `Rstn`.

7. If you are compiling the HyperTransport MegaCore function variation file top-level entity in your Quartus II project, set virtual pin attributes for all of the internal interface signals of the variation.

   ☞   An example Quartus II project that has all of the above I/O standards set, and virtual pin and clock latency settings, is included with the HyperTransport MegaCore function installation. Refer to "Example Quartus II Project" on page 2–16.

8. Turn on the Quartus II timing analysis setting **Enable Clock Latency** to perform correct timing analysis. Refer to Quartus II Help for instructions on how to make this assignment.

9. Set the remaining constraints in the Quartus II software, including the device, pin assignments, timing requirements, and any other relevant constraints. Refer to Appendix B, Stratix Device Pin Assignments for more details about assigning pins. If you have not made pin assignments for your board design, you can use the Quartus II software to automatically assign pins.

10. On the Processing menu, click **Start Compilation** to compile the design.

Program a Device

After you compile your design, you can program your targeted Altera device and verify your design in hardware.

With Altera's free OpenCore Plus evaluation feature, you can evaluate the HyperTransport MegaCore function before you purchase a license. OpenCore Plus evaluation allows you to generate an IP functional simulation model and produce a time-limited programming file.

You can simulate the HyperTransport MegaCore function in your design and perform a time-limited evaluation of your design in hardware.

☞   For more information about OpenCore Plus hardware evaluation for the HyperTransport MegaCore function, refer to "OpenCore Plus Time-Out Behavior" on page 3–40 and *AN320: OpenCore Plus Evaluation of Megafunctions*.

Set Up Licensing

You need purchase a license for the MegaCore function only after you are completely satisfied with its functionality and performance and want to take your design to production.

After you purchase a license for the HyperTransport MegaCore function, you can request a license file from the Altera web site at www.altera.com/licensing and install it on your computer. After you request a license file, Altera emails you a **license.dat** file. If you do not have Internet access, contact your local Altera representative.

To install your license, you can either append the license to your Quartus II software **license.dat** file or you can specify the MegaCore function's **license.dat** file in the Quartus II software.

☞ Before you set up licensing for the HyperTransport MegaCore function, you must already have the Quartus II software installed on your computer with licensing set up.

Append the License to Your license.dat File

To append the license, follow these steps:

1. Close the following software if it is running on your computer:

   ■ Quartus II software

   ■ MAX+PLUS® II software

   ■ LeonardoSpectrum™ synthesis tool

   ■ Synplify software

   ■ ModelSim® simulator

2. Open the HyperTransport license file in a text editor. The file should contain one `FEATURE` line, spanning 2 lines.

3. Open your Quartus II **license.dat** file in a text editor.

4. Copy the `FEATURE` line from the HyperTransport license file and paste it in the Quartus II license file.

   ☞ Do not delete any `FEATURE` lines from the Quartus II license file.

5. Save the Quartus II license file.

   ☞ When using editors such as Microsoft Word or Notepad, ensure that the file does not have extra extensions appended to it after you save (for example, **license.dat.txt** or **license.dat.doc**). Verify the file name in a DOS box or at a command prompt.

Specify the License File in the Quartus II Software

To specify the MegaCore function's license file, follow these steps:

☞ Altera recommends that you give the file a unique name, for example,
    *<MegaCore name>*_**license.dat**.

1. Run the Quartus II software.

2. On the Tools menu, click **License Setup**. The **Options** dialog box opens to the
    License Setup page.

3. In the **License file** box, add a semicolon to the end of the existing license path and
    file name.

4. Type the path and file name of the MegaCore function license file after the
    semicolon.

    ☞ Do not include any spaces either around the semicolon or in the path or file
        name.

5. Click **OK** to save your changes.

Example Simulation and Compilation

Altera provides example design files in the directory *<path>***ht****example**, where
<path> is the directory in which you installed the MegaCore function. You can use
these design files to run vector-based simulations and to compile in the Quartus II
software. The example can help you validate the installation of the HyperTransport
MegaCore function in your design environment and serve as an example for setting
up your custom environment.

Example Quartus II Project

Altera provides an example Quartus II project in the directory
<path>**ht****example****quartus** that compiles the file *<path>***ht****example****ht_top.v**.
This project has Quartus II virtual pins assigned for the user-side signals from the
MegaCore function variation. The target device is the EP1S60F1020C6. The pin
assignments used for the HyperTransport link signals support all clocking options.

To compile this example Quartus II project, perform the following steps in the
Quartus II software version 9.0:

1. Start the Quartus II software.

2. On the File menu, click **Open Project**.

3. In the **Open Project** dialog box, browse to *<path>*\ **ht****example****quartus**.

4. Select the **ht_top.qpf** file.

5. On the Processing menu, click **Start Compilation** to compile the project.

Example Simulation with Test Vectors

The example design in directory *<path>***ht****example** contains the following Verilog
HDL files used in the example simulation:

■ **ht_top.vo**

■ **ht_top_tb.v**

☞ The file **ht_top.vo** is the IP functional simulation model corresponding to the **ht_top.v** variation file included the this directory.

The test vectors consist of the following files:

■ **input_ht_vector.dat**

■ **output_ht_vector.dat**

■ **input_ui_vector.dat**

■ **output_ui_vector.dat**

When you simulate your design, the top module in the simulation file **ht_top_tb.v** reads the input vectors at startup and drives them as test vectors to **ht_top.vo**. It compares the outputs from **ht_top.vo** to the output vectors. If the simulation module detects a mismatch, it displays an error and the simulation stops. If all of the vectors match, the simulation displays a message indicating that the test passed. The test vectors are described in "About the Test Vectors" on page 2–18.

☞ The following simulation examples use an IP functional simulation model generated with the Quartus II software version 9.0. When simulating your design, you must use the simulation library files supplied with the version of the Quartus II software you used to create the IP functional simulation model for your MegaCore function variation.

Simulating Test Vectors Using ModelSim

To simulate with the test vectors interactively, perform the following steps:

1. Start the ModelSim software by executing the **vsim** command at a system command prompt or by using the Windows Start menu.

2. In the Modelsim session, change your working directory to the Modelsim example project directory *<path>*/**ht/example/modelsim**.

3. To map the correct libraries, compile the required files, and load and run the simulation, type the following command:

    do run.do ↵

To simulate with the test vectors without opening a Modelsim session, perform the following steps:

1. Change your working directory to the ModelSim example project directory *<path>*/**ht/example/modelsim**.

2. To map the correct libraries, compile the required files, and load and run the simulation, perform one of the following steps:

    ■ On a supported Windows platform, double-click **run.bat** or, if scripting, type the following command:

        run ↵

    ■ On a supported Linux operating system, at a system command prompt, type the following command:

        ./run.sh ↵

About the Test Vectors

The example design has a host connected to an instance of the HyperTransport MegaCore function configured as an end-chain link. When simulating the receive link interface, the following events occur:

1. The host configures the end chain link's `BAR0` with `0xfa000000`. It then writes `0x01` as the `UnitID` and `0x0007` in the command register of the device header space.

2. The host writes double-word data with data ranging from 1 to 16 DWORDs to the posted channel followed by a byte write with data ranging from 1 to 8 DWORDs.

3. The host writes double-word data with data ranging from 1 to 16 DWORDs to the non-posted channel followed by a byte write with data ranging from 1 to 8 DWORDs.

4. The host performs a read request with data ranging from 1 to 16 DWORDs to the non-posted channel followed by a byte read request with data ranging from 1 to 8 DWORDs.

Analyzing the transmit local-side interface, the following sequence of events occurs:

1. The local side writes double word data with data ranging from 1 to 16 DWORDs to the posted channel followed by a byte write with data ranging from 1 to 8 DWORDs.

2. The local side writes double word data with data ranging from 1 to 16 DWORDs to the non-posted channel followed by a byte write with data ranging from 1 to 8 DWORDs.

3. The local side performs a read request with data ranging from 1 to 16 DWORDs to the non-posted channel followed by a byte read request with data ranging from 1 to 8 DWORDs.

HyperTransport Technology Overview

HyperTransport technology (HT) is packet-based point-to-point link that is designed to deliver a scalable, high-performance interconnect between the CPU, memory, and I/O devices on a circuit board. The HT link uses low-swing differential signaling with differential termination to achieve high data rates from 400 Megabytes per second (Mbytes/s) to 1.6 Gigabytes per second (Gbytes/s) per direction, assuming an 8-bit interface.

The HT link provides significantly more bandwidth than competing interconnect technologies; it uses scalable frequency and data width to achieve scalable bandwidth. Designers can use HT in networking, telecommunications, computer, and high-performance embedded applications, and in applications that require high speed, low latency, and scalability.

The HT link consists of two independent, source-synchronous, clocked, and unidirectional sets of wires, as illustrated in Figure 3–1.

Figure 3–1. HT Link
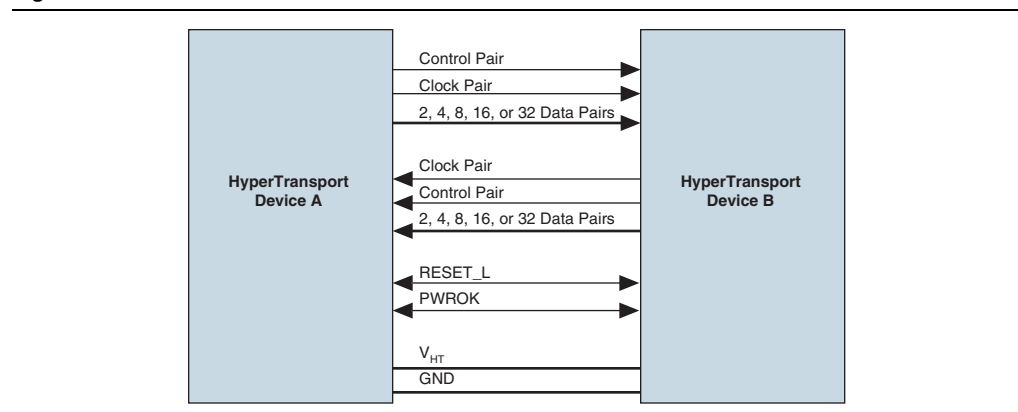
For additional information, refer to the *HyperTransport I/O Link Specification Revision 1.03*.

HT systems consist of two or more HT devices. The HyperTransport specification includes the following device types:

■ *Host Bridge*—A host bridge is the HT interface that provides connectivity to the system's host processor. Because all communication within an HT chain is between individual devices and the host bridge, the host bridge includes additional functionality such as managing peer-to-peer packets and handling error conditions.

■ *End-Chain Link*—The end-chain link device is the simplest of all HT devices and only one can exist in an HT chain. This single-link device claims packets when the destination address of the packet matches its address. If it receives a packet that does not match its address, it must indicate an error condition by setting appropriate error bits or sending a return packet indicating that an error occurred. The Altera HyperTransport MegaCore function implements an end-chain link.

■ *Tunnel*—A tunnel is a dual-link device that is not a HyperTransport-to-HyperTransport bridge. In the upstream direction, tunnels transfer all packets except for information packets directly from the downstream link to the upstream link. In the downstream direction, only packets that are not claimed by the device are transferred to the downstream interface.

■ *HyperTransport-to-HyperTransport Bridge*—This bridge has one or more HT links and is more complex than a tunnel because it translates packets from one chain to two or more chains. The bridge must take upstream traffic from secondary chains and forward the packets to the primary chain while maintaining I/O streams, virtual channel ordering, and error conditions.

HT Systems

HT systems, or fabrics, are implemented as one or more daisy chains of HT devices. Figure 3–2 shows two of the three variations of HT fabric configurations. You can use the HyperTransport MegaCore function as the end-chain link in any of these systems.

Figure 3–2. HT Systems
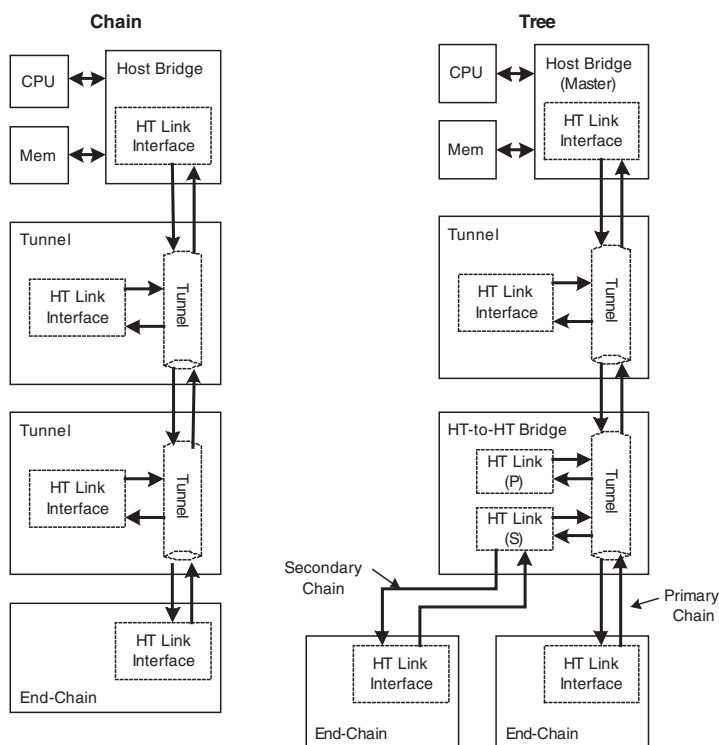
HT Flow Control

All commands and data are separated into one of three separate virtual channels, which provide higher performance and allow certain types of requests to pass others to avoid deadlocks. However, in some cases, requests in one channel cannot pass those in other channels to preserve ordering as required by systems. The three virtual channels are:

■ *Non-Posted Requests*—Requests that require a response (all read requests and optionally write requests)

■ *Posted Requests*—Requests that do not require a response (typically write requests)

■ *Responses*—Responses to non-posted requests (read responses or target done responses to non-posted writes)

The HT flow control mechanism is a credit-based scheme maintained per virtual channel for an individual link. A transmitter consumes a buffer credit each time it transmits a packet and cannot transmit a packet to the receiver unless a buffer credit is available. The receiver provides buffer credits for each available buffer at link initialization, and it provides an additional buffer credit each time a buffer is freed thereafter.

Buffer credits are transmitted in the opposite direction of the data flow as part of NOP packets.

HyperTransport MegaCore Function Specification

This section describes the functionality and features of the 8-bit end-chain HyperTransport MegaCore function. Figure 3–3 shows the block diagram of the HyperTransport MegaCore function. The HyperTransport MegaCore function is partitioned into three layers:

■ Physical Interface

■ Synchronization and Alignment

■ Protocol Interface

Figure 3–3. HyperTransport MegaCore Function Block Diagram
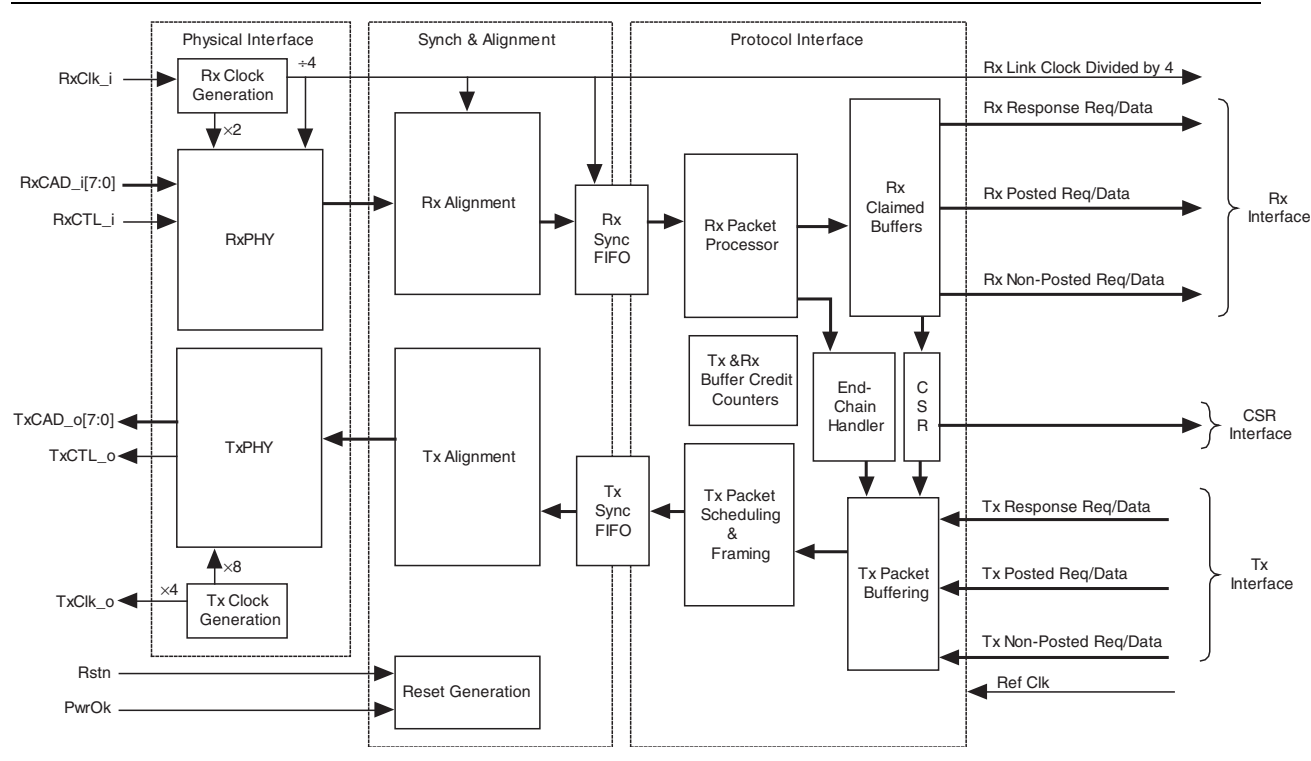
Physical Interface

The physical interface module contains the logic that interfaces the HyperTransport MegaCore function to the physical link signals. It contains the logic for both the receiver and transmitter portions. The SERDES functionality is embedded into the Altera device to ensure operation at maximum speed.

Rx Deserialization Logic and Clock Generation

This module receives serial data from each `RxCTL_i` and `RxCAD_i` bit and converts it to a parallel data stream. The Rx clock generator (implemented in a fast PLL) multiplies the receiver Rx link `RxClk_i` by two to capture `RxCTL_i` and `RxCAD_i`. The data is then deserialized by a factor of eight and clocked by a divided-by-four clock for use in the synchronization and alignment layer.

Tx Serialization Logic and Clock Generation

This module transmits serial data on the link interface. Data is transferred from the internal module in parallel and is serialized on the link interface. The data entering this module is already split into channels. The serialization circuit drives `TxCAD_o` and `TxCTL_o`. The Tx clock generator multiplies the Tx alignment logic clock by eight to serialize the `TxCTL_o` and `TxCAD_o` signals, and generates a multiplied-by-four clock for use as the HT `TxClk_o`.

Synchronization and Alignment

The synchronization and alignment layer is responsible for synchronizing clock domains and aligning data to the natural boundaries.

Rx Synchronization and Alignment Interface

The Rx synchronization and alignment interface performs the following functions:

■ *Bit-to-Byte Alignment*—The interface performs link initialization sync packet detection and byte alignment so that the first received byte of data is placed in byte 0 of the internal 64-bit data path.

■ *CRC Checking*—The interface checks the link CRC.

■ *Synchronization*—The interface writes all non-CRC data to an Rx synchronization FIFO buffer so that the data can be synchronized to the protocol interface clock domain for the protocol interface layer.

  ☞ If you turn on the **Shared Rx/Tx/Ref Clock** option in the IP Toolbench parameterization wizard, the IP Toolbench removes the Rx synchronization FIFO buffer from the design to reduce latency. Refer to "Clocking Options" on page 3–7 for more information about clock options.

Tx Alignment Interface

The Tx alignment interface performs the following functions:

■ *Synchronization*—The interface reads transmit data from the Tx synchronization FIFO buffer to move the data from the protocol interface clock domain to the Tx alignment clock domain.

  ☞ If you turn on the **Shared Ref/Tx Clock** or **Shared Rx/Tx/Ref Clock** option in the wizard, the wizard removes the Tx synchronization FIFO buffer from the design to reduce latency. Refer to "Clocking Options" on page 3–7 for more information about clock options.

■ *Link Initialization*—The interface generates the link initialization sequence when the link is reset.

■ *CRC Generation*—The interface generates CRCs.

Protocol Interface

The protocol interface module has a 64-bit data path. It receives formatted packets from the Rx synchronization and alignment module and transmits packets to the Tx alignment module. Because the HT specification requires that traffic in the three different virtual channels (posted, non-posted, and responses) be kept independent, the module maintains internal packet buffering and the local interface separately for each virtual channel. The following sections describe the protocol interface blocks and their function.

Rx Packet Processor

The Rx packet processor reads the data stream from the Rx sync FIFO buffer. It parses the data and determines whether packets should be claimed or passed to the end chain handler.

If the Rx request address matches one of the BAR registers in the CSR module, the Rx packet processor claims the packet and writes to the Rx posted buffer or the Rx non-posted buffer. If the UnitID in an Rx response matches the UnitID in the CSR, the processor claims the packet and writes to the Rx response buffer. The processor claims data packets and writes to the appropriate Rx buffer if it claimed the associated request or response packet.

The Rx packet processor passes unclaimed packets to the end-chain handler.

End-Chain Handler

The end-chain handler logs errors when it receives a response or posted packet. It also generates NXA response packets when it receives a non-posted packet.

Rx Claimed Buffers

The traffic written to the claimed packet buffers is stored in the appropriate virtual channel buffer. The packets are stored in the order in which they are received from the link. When reading the packets from the buffers, you read commands followed by any associated data.

To allow maximum throughput of the command packets, command packets are stored in registers. Data packets, on the other hand, are stored in dual-port memory blocks.

Although data packets and command packets are stored in separate storage elements, when the user interface reads those packets it appears as though they are stored in the same location. That is, the user interface sees a command packet followed by the data, consecutively.

Tx Buffers

The Tx buffers are temporary storage for the traffic to be transmitted on the Tx path. The user logic writes packets to the buffer with the command as the first word written followed by any data associated with that command. Tx buffers automatically adjust for 32-bit commands or 64-bit commands by inserting idle NOP packets in the upper bits of a 32-bit command, such as a read response. Additionally, because the transmitted data may be an odd number of DWORDS, the Tx buffers insert a `NOP` packet from the NOP generator to align the end of packet to the 64-bit boundary, placing the next packet start command on the lower DWORD. The Tx buffers also generate appropriate CTL information for transmission to the link.

Scheduler

The scheduler ensures equal access to all HT virtual channels. The scheduler is an arbiter that performs round-robin arbitration between the response, non-posted, and posted buffers. Additionally, the scheduler gives a higher priority to the NOP generator so that if a NOP packet with flow control information is available at the end of a packet transmission, the scheduler allows it to be transmitted before starting a new packet transmission from another virtual channel.

Figure 3–4 shows example transmitter traffic. This example assumes that each square represents one DWORD. The top square is the low DWORD and the bottom square is the high DWORD. In Figure 3–4, the leftmost information transmits first. The vertical dashed lines delineate packets. The idle NOP packets are NOP packets inserted to align data and never contain flow control information. Flow control NOP packets are read from the NOP generator and may or may not contain flow control information, depending on the buffer release status generated by the Rx claimed buffers.

Figure 3–4. Example Transmission in the Tx Interface

CSR Module

The CSR module reads packets from the non-posted buffer and checks whether they are CSR access packets. This process allows the HyperTransport MegaCore function to use the same buffers for the CSR interface that it uses for general non-posted traffic, reducing the size of the MegaCore function variation. If a packet is a CSR access (read or write request) it is routed to the CSR interface module instead of the user interface. Normal non-posted traffic is provided directly to the local user interface.

The CSR module contains all of the CSR registers and provides read/write capability to them. It also generates the appropriate response to CSR accesses in the Tx response buffer.

Clocking Options

The HyperTransport MegaCore function has three distinct clock domains:

- Protocol interface clock domain

- Rx alignment clock domain

- Tx alignment clock domain

The MegaCore function has three clocking options that link clock domains together to reduce overall latency and resource usage. Due to the HT protocol flow control mechanism, the clocking option you choose results in different latency for the flow control information inside the HyperTransport MegaCore function, which can affect your overall system performance.

Shared Rx/Tx Clock

When you turn on the **Shared Rx/Tx Clock** option, the Rx alignment logic, protocol interface logic, and Tx alignment logic operate as independent clock domains with synchronization FIFO buffers buffering data between the modules, as shown in Figure 3–5. The Rx and Tx clocks share the same PLL, and are synchronous. This implementation provides the most flexible design because the reference clock (`RefClk`) frequency and phase can be determined by other requirements the user application may impose, and the HT Tx link clock is always synchronous with the HT Rx link clock.

Figure 3–5. Shared Rx/Tx Clock Option

The Tx sync and Rx sync FIFO buffers consume logic resources and add latency to the HyperTransport MegaCore function. This increased latency can limit the link throughput if the attached HT device has a small number of buffer credits. The latency increases the turnaround time for receiving a new buffer credit after transmitting a packet.

The design must meet the following clock frequency requirements when using the **Shared Rx/Tx Clock** option:

- The frequency of `RefClk` must be greater than or equal to `RxLnkClkD4`.

- When `RefClk` and `RxLnkClkD4` are nominally equal but are derived from different sources, `RefClk` must be no more than 2,000 ppm slower than `RxLnkClkD4`.

☞ Failing to meet these requirements will result in system failure due to Tx sync FIFO buffer underflow or Rx sync FIFO buffer overflow.

In most designs, you would not connect `RxLnkClkD4` to anything external to the HyperTransport MegaCore function, but it is provided for monitoring purposes if needed.

Shared Ref/Tx Clock

The **Shared Ref/Tx Clock** option reduces the latency through the HyperTransport MegaCore function by eliminating the Tx sync FIFO buffer. The MegaCore function uses `RxLnkClkD4` for the Rx alignment logic only, and uses the `RefClk` for the rest of the logic, including the Tx alignment logic, as shown in Figure 3–6.

Figure 3–6. Shared Ref/Tx Clock Option
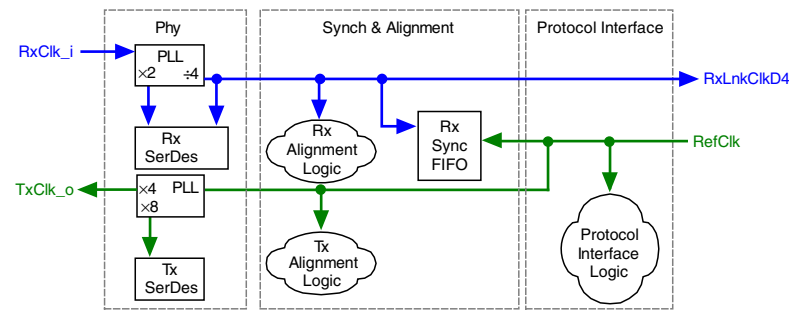
The design must meet the following clock frequency requirements when using the **Shared Ref/Tx Clock** option:

■ The frequency of `RefClk` must be greater than or equal to `RxLnkClkD4`.

■ When `RefClk` and `RxLnkClkD4` are nominally equal but are derived from different `RxLnkClkD4`, RefClk must be no more than 2,000 ppm slower than RxLnkClkD4.

■ `RefClk` should run at 50, 75, 100, or 125 MHz to create the allowed HT clock frequencies of 200, 300, 400, or 500 MHz.

☞ Failing to meet these requirements results in system failure due to Tx sync FIFO buffer underflow or Rx sync FIFO buffer overflow.

Additionally, the attached HT device may require its incoming HT link clock to be less than or equal to its outgoing HT link clock. This requirement—along with the above requirements—forces the `RefClk` frequency to be within 2,000 ppm of the `RxLnkClkD4` frequency.

☞ This clocking option may not be strictly compliant with the HT specification. If `RefClk` runs at a frequency other than 50 MHz, `TxLnkClk` does not run at the required 200 MHz upon reset. However, in many embedded applications it may be acceptable for the `TxLnkClk` to operate at 300, 400, or 500 MHz upon reset.

Depending on the time base of the supplied `RefClk`, the **Shared Ref/Tx Clock** option implements an asynchronous or synchronous HT implementation. The time base for the `RefClk` can be asynchronous or synchronous with the attached receiver's time base.

In most designs, you would not connect `RxLnkClkD4` to anything external to the HyperTransport MegaCore function, but it is provided for monitoring purposes if needed.

Shared Rx/Tx/Ref Clock

When you use the **Shared Rx/Tx/Ref Clock** option, `RxLnkClkD4` is used for the entire HyperTransport MegaCore function and the local side interfaces. This implementation provides the lowest latency MegaCore function variation and typically results in the highest performance, and is a synchronous HT clock implementation. All of the user logic that interfaces to the HyperTransport MegaCore function must operate using `RxLnkClkD4`. This option is illustrated in Figure 3–7.

Figure 3–7. Shared Rx/Tx/Ref Clock Option
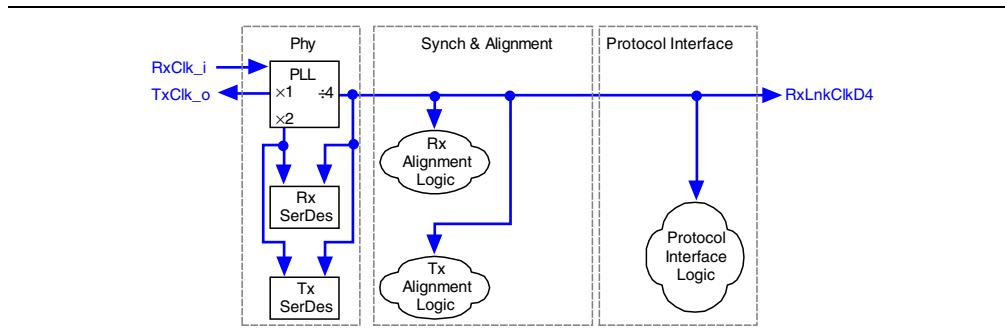
In this case, the Rx sync FIFO buffer and the Tx FIFO buffer are removed, yielding the best throughput performance (i.e., the lowest latency through the MegaCore function variation.)

☞ You cannot use the `RefClk` input if you use the **Shared Rx/Tx/Ref Clock** option.

HyperTransport MegaCore Function Parameters and HT Link Performance

This section describes how the Rx buffer size parameters and clocking options of the HyperTransport MegaCore function relate to throughput on the HT Link interface. Refer to section 4.8 of the *HyperTransport I/O Link Specification Revision 1.03* for information about the flow control mechanism before reading this section.

The HT flow control mechanism is a credit-based scheme in which the transmitter maintains a counter for each type of buffer in the receiver. When the link initializes, the transmitter resets all of the counters to zero. When link initialization completes, each receiver has its transmitter send buffer credits (transmitted within NOP commands) for each type of buffer to indicate the number of buffers available for each type of packet. When the transmitter transmits a packet of a particular type it decrements that counter by one. The transmitter cannot transmit a packet if the count for that type of packet is 0. After the receiver processes the packet and frees up the buffer, it has its transmitter send a buffer credit for that type back to the original transmitter.

The transmitter maintains six counters, one for each type of packet. If a command packet has an associated data packet, the transmitter must ensure that it has credits for both the control and data buffer types. Therefore, for most traffic on the HT link interface, think of a pair of data and control counters as one. For simplicity, the rest of this discussion refers to the control and data counters for each virtual channel as one counter.

For applications that require large bursts of data to be transmitted on one virtual channel, the counter of that virtual channel can limit the throughput. This situation occurs because when the counter reaches zero, no additional packets can be transmitted in that virtual channel. If another virtual channel has no packets to be transmitted, the link is idle until more credits are received. To maximize the throughput in this type of application, you must prevent or minimize idle time on the link by ensuring that new credits are received before the counter reaches zero. This implementation ensures continuous transfer until all data is transmitted.

Figure 3–8 shows the components of the HT flow control loop for a single virtual channel. The loop is from the time Transmitter A decrements its available Rx buffer counter until the time that counter is incremented.

Figure 3–8. HT Flow Control Loop *(Note 1)*
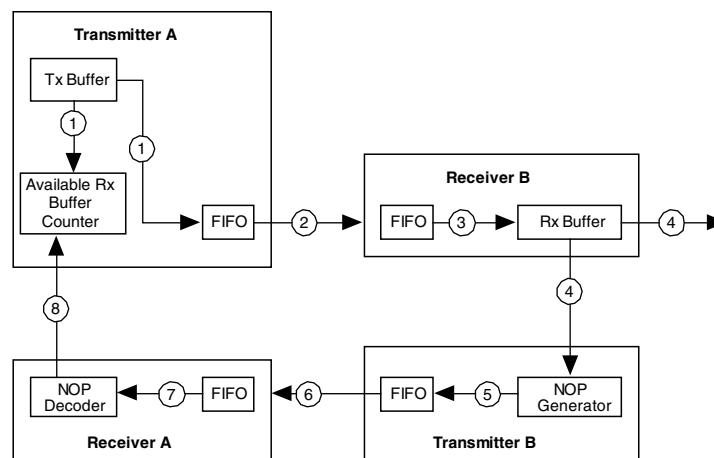
Note to Figure 3–8:

(1)   Numbered steps are described in the following sections.

Figure 3–8 demonstrates the following flow:

1. Transmitter A schedules a packet for transmission and transfers it to the Tx FIFO buffer. At the same time, the locally maintained available Rx buffer counter is decremented.

2. The command is read out of the Tx FIFO buffer and transferred across the HT link to receiver B's Rx FIFO buffer.

3. The command is read out of the Rx FIFO buffer, decoded, and written into the Rx buffer.

4. The user-side logic reads the command from the Rx buffer and frees the buffer. The buffer free indication is transferred to the NOP generator in transmitter B.

5. The NOP containing the Rx buffer credit information is scheduled for transmission and written into transmitter B's Tx FIFO buffer.

6. The NOP is read out of transmitter B's Tx FIFO buffer and transferred across the HT link to receiver A's Rx FIFO buffer.

7. The NOP is read out of receiver A's Rx FIFO buffer and decoded.

8. The Rx buffer credit is sent to transmitter A's available Rx buffer counter, allowing it to be incremented.

If the number of Rx buffers is large enough, the virtual channel can have a high loop time and continuous transmission. For example, if you assume that transmitter A has not transmitted packets for a particular virtual channel for some time, the avaliable Rx buffer counter in transmitter A is equal to the total number of Rx buffers for that virtual channel in receiver B. If transmitter A now begins transmitting a large burst of packets on that virtual channel, the first buffer credit must make it back to increment the counter before it has decremented to zero to ensure continuous transmission without interruption. This process repeats and allows the burst to be transmitted uninterrupted by the counter until the burst is finished.

On the other hand, if the number of Rx buffers is small and the flow control loop time is high, the counter decrements to zero before the first buffer credit makes it back. In this case, the burst is interrupted by the counter until the first buffer credit makes it back. Assuming no other virtual channels have packets available to transmit, transmitter A is forced to transmit NOPs, limiting the overall utilization of the link. After the first buffer credit is received, transmitter A immediately transmits another packet and decrements the counter to zero. The virtual channel is stalled until the next buffer credit is received.

The typical packet size also affects the throughput. Larger packets take longer to transmit, so the decrement rate of the available Rx buffer counter is relatively slow. Thus, large packets permit a longer loop time because the counter takes longer to decrement to zero.

☞ Non-posted reads are command-only requests. A burst of non-posted reads can be transmitted very rapidly and a large number of available non-posted Rx buffers can be exhausted quickly. However, the overall throughput of reads is actually limited by the throughput of the response virtual channel in the opposite direction where both command and data packets are being transferred. As long as the number of non-posted Rx buffers is greater than or equal to the number of Rx response buffers available at the opposite end of the link, the Rx non-posted buffer count does not limit the overall read throughput.

In applications that use a mix of packets on multiple virtual channels, the decrement rate of the available Rx buffer counter on a particular channel is slower. This rate allows a smaller number of Rx buffers per virtual channel or a longer loop time.

The HyperTransport MegaCore function provides two features that allow you to adjust actual throughput, depending on the HT device to which the MegaCore function is connected:

■ Rx buffer size

■ Clocking options

The HyperTransport MegaCore function allows you to adjust the Rx buffer size for each virtual channel. Therefore, reasonably long flow control loop times and typical packet sizes of 32 bytes, as opposed to the maximum 64 bytes, can be tolerated with no loss in throughput for traffic received by the HyperTransport MegaCore function.

When the attached HT device has a limited number of Rx buffers (on the order of four or less), you can adjust the HyperTransport MegaCore function clocking option to minimize the loop time. In this case, the flow control loop time is critical for maintaining high throughput on HT link traffic transmitted from the HyperTransport MegaCore function. By adjusting the clocking option, you can remove either the Tx FIFO buffer or both the Tx and Rx FIFO buffers. Removing one or both buffers increases throughput by removing the latency through the FIFO buffers. If you must maintain maximum throughput on a single virtual channel with a small number of Rx buffers, then use either the **Shared Ref/Tx Clock** or **Shared Rx/Tx/Ref Clock** option. Using the **Shared Ref/Tx Clock** option removes the Tx FIFO buffer. Using the **Shared Rx/Tx/Ref Clock** option removes both the Tx and Rx FIFO buffers.

Refer to Table A–3 on page A–3 and Table A–4 on page A–3 for information about the clocking option and Rx buffer size parameters.

Signals

Figure 3–9 shows a top-level view of the HyperTransport MegaCore function interface signals.

Figure 3–9. Top-Level HyperTransport MegaCore Function Sign
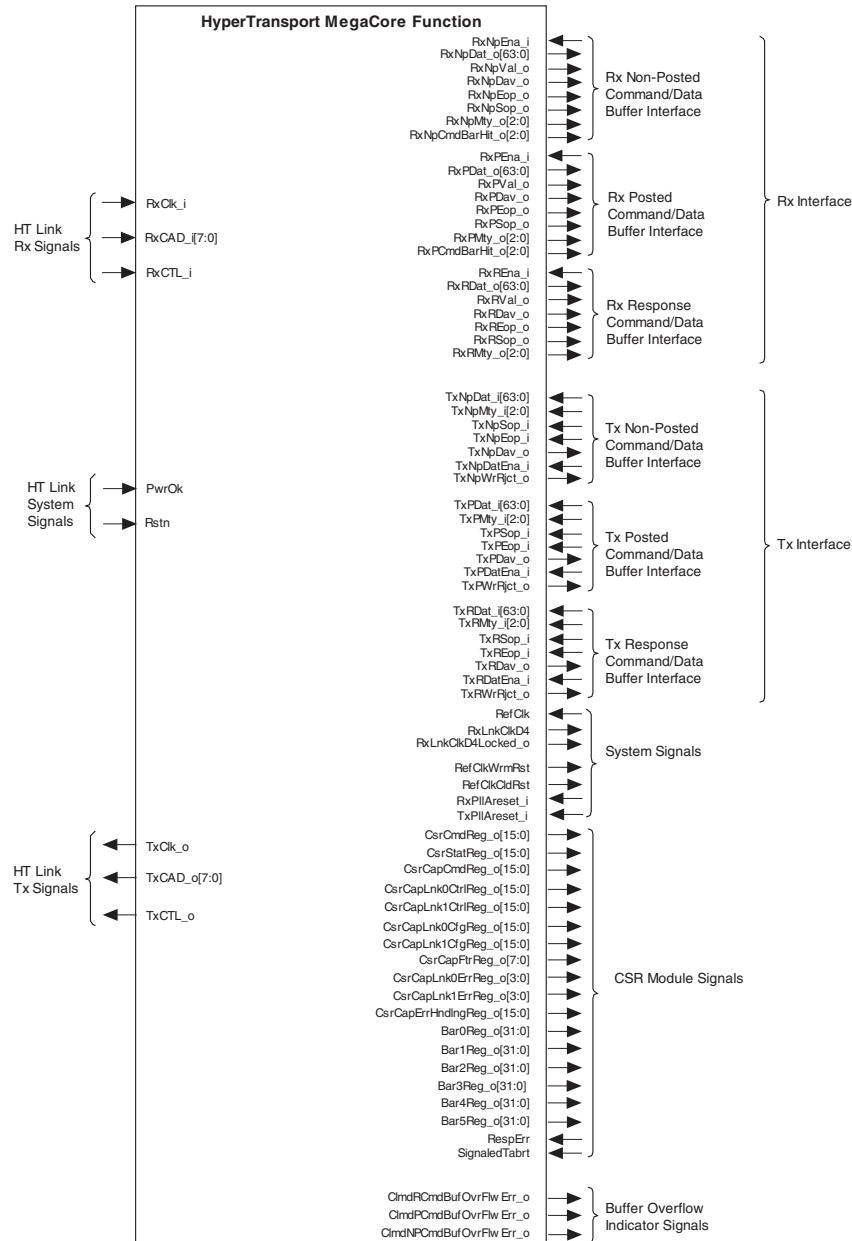
The HT interface signals are also external I/O pins on the chip. Each HT interface signal is in one of the following groups:

- HT link Rx signals
- HT link system signals
- HT link Tx signals

The local application interface uses an Altera Atlantic interface to transfer data to the local-side application. The local application interface has the following signal groups:

■ Rx non-posted command/data interface

■ Rx posted command/data interface

■ Rx response command/data interface

■ Tx non-posted command/data interface

■ Tx posted command/data interface

■ Tx response command/data interface

■ CSR interface

■ System signals

The following sections describe these signal groups.

The local application interfaces to the Rx and Tx command/data buffers using the Altera Atlantic interface. For a detailed description of the Atlantic interface, refer to *FS 13: Atlantic Interface*.

Table 3–1 offers an overview of the Atlantic interface information in *FS 13: Atlantic Interface* that is relevant for the HyperTransport MegaCore.

Table 3–1. Atlantic Interface Support

Interface	Atlantic Type	Unused Signals	Exceptions
Rx non-posted command/data interface. Rx posted command/data interface. Rx response command/data interface.	Slave Source	`err` `par` `adr`	Packets are organized in the form of command followed by data. The Dat signals hold command and data information. When SOP is asserted, the information driven on the Dat signal is command information. Side-band signals provide additional information.
Tx non-posted command/data buffer signals Tx posted command/data signals. Tx response command/data signals	Slave Sink	`err` `par` `adr`	The Rjct signal indicates when the write was rejected. This signal is a side-band signal in addition to the Atlantic signals.
CSR signals	Not Atlantic		
System signals	Not Atlantic		

The Atlantic interface specification defines the `dat` signal bus to be in big-endian format. Because HT uses little-endian format, the HyperTransport MegaCore function uses a little-endian format across the Atlantic interface. This implementation requires a different interpretation of the Atlantic `mty` signals, as described in the `mty` signal descriptions in the remainder of this section.

☞ The following description distinguishes between an HT packet and an Atlantic packet. An HT packet is defined as indicated in the HT specification with various types including 32-bit commands, 64-bit commands, or various size data packets. An Atlantic packet, however, contains one or two HT packets. An Atlantic packet can contain only a single HT command packet (32- or 64-bit), or a single HT command packet followed by its associated HT data packet.

HT Link Rx Signals

The HT link Rx signals are differential pair signals. Table 3–2 defines a single signal name for each pair, and the same is true of the HyperTransport MegaCore port declarations. These signals must be assigned the HyperTransport I/O standard in the Quartus II software; the software creates the differential pair automatically during compilation.

Table 3–2. HT Link Rx Signals

Signal Name	Direction	Description
RxClk_i	Input	RxClk_i is the received HT clock and is used to clock the RxCAD_i and RxCTL_i inputs.
RxCAD_i[7:0]	Input	RxCAD_i[7:0] is the received HT command, address, and data bus.
RxCTL_i	Input	RxCTL_i is the received HT control signal.

HT Link System Signals

Table 3–3 describes the HT link system signals. These signals are 2.5-V tolerant LVCMOS. The HyperTransport MegaCore function does not implement the optional HT power management signals, LDTSTOP# and LDTREQ#.

Table 3–3. HT Link System Signals

Signal Name	Direction	Description
Rstn	Input	Rstn is the HT Reset# signal. HT allows it to be a wired OR signal with multiple drivers, however, the HyperTransport MegaCore function treats it as input only.
PwrOk	Input	PwrOk is the HT power okay signal. HT allows it to be a wired OR signal with multiple drivers, however, the HyperTransport MegaCore treats it as input only.

HT Link Tx Signals

The HT link Tx signals are differential pair signals. Table 3–4 defines a single signal name for each pair, and the same is true of the HyperTransport MegaCore port declarations. These signals must be assigned the HyperTransport I/O standard in the Quartus II software; the software creates the differential pair automatically during compilation.

Table 3–4. HT Link Tx Signals

Signal Name	Direction	Description
TxClk_o	Output	TxClk_o is the transmitted HT clock and is center-aligned with respect to the TxCAD_o[7:0] and TxCTL_o signals.
TxCAD_o[7:0]	Output	TxCAD_o[7:0] is the transmitted HT command, address, and data bus.
TxCTL_o	Output	TxCTL_o is the transmitted HT control signal.

Rx Command/Data Buffer Interfaces

Each HT virtual channel has one Rx command/data interface. These interfaces are slave source Atlantic interfaces, allowing streaming packets in an easy-to-use interface.

Rx Command/Data Buffer Interface Signals

Table 3–5 describes the Rx buffer interface signals. Each of these signals has one of the following three types:

■ Rx response interface (HT MegaCore function signal names prefixed with RxR)

■ Rx posted interface (HT MegaCore function signal names prefixed with RxP)

■ Rx non-posted interface (HT MegaCore function signal names prefixed with RxNp)

Because each channel uses similar interface signals, they are only described once. Any channels that have specific differences are noted in Table 3–5.

Table 3–5. Rx Command/Data Buffer Interface Signals (Part 1 of 2)

Signal Name	Direction	Description
Dat_o[63:0]	Output	Data bus. This 64-bit bus contains the actual HT command/data information received from the HT link. This bus has the following constraints: ■ Data received from the link is organized in Atlantic packets. ■ Data is transferred in a little endian format (first byte is in Dat_o[7:0]). ■ The Atlantic packet on the Rx interface contains a command phase followed by zero, one, or more data phases. ■ During the command phase, Sop_o is asserted to indicate the start of packet. That is, Sop indicates the start of an Atlantic packet and indicates that Dat_o[63:0] contains the HT command. ■ The data presented on Dat_o[63:0] is presented with the bytes in the order they are received from the HT. Therefore, the first byte received from the HT link in a command packet is provided on Dat_o[7:0]. ■ If the command has no associated data, such as a non-posted read request or TgtDone response packet, Sop_o and Eop_o are asserted at the same time. ■ You can decode the command bits to determine whether there is data associated with the command (and how much) while Sop_o is asserted.
Mty_o[2:0]	Output	Data byte empty. This bus indicates which bytes are invalid on Dat_o[63:0]. Because HT packets are DWORD aligned, only two valid encodings are used by the HyperTransport MegaCore function. The following Mty_o signal values are possible: ■ Mty_o = '000', all bytes on Dat_o[63:0] are valid. ■ Mty_o = '100', Dat_o[63:32] are invalid. Invalid bytes can only be on the last word of a packet. It is illegal to have a non-zero value for the Mty_o signals for words other than the last word of the packet, even for 32-bit commands that have data, such as a read response. In this case, when Sop_o is asserted, Dat_o[63:32] is implicitly invalid. The first data bytes are placed on the Dat_o bus the cycle after Sop_i. If the 32-bit command does not have data, the Mty_o signals are set to '100', and Eop_o and Sop_o are asserted.

Table 3–5. Rx Command/Data Buffer Interface Signals (Part 2 of 2)

Signal Name	Direction	Description
Sop_o	Output	Start of packet. This signal indicates the start of packet. When Sop_o is high, start of packet is present on Dat_o[63:0] and is aligned to the least significant byte. Sop_o is qualified with the Val_o signal. If Val_o is low, Sop_o must be ignored.
Eop_o	Output	End of packet. This signal indicates the end of packet. When Eop_o is high, end of data packet is present on Dat_o[63:0] and is aligned to the least significant byte. Eop_o is qualified with the Val_o signal. If Val_o is low, Eop_o must be ignored.
Val_o	Output	Data valid. This signal indicates that the data driven on Dat_o[63:0] is valid. Val_o is updated on every RefClk edge at which Ena_i is sampled asserted, and holds its current value along with the Dat_o bus when Ena_i is sampled deasserted. When Val_o is asserted, the Atlantic data interface signals are valid. When Val_o is deasserted, the Atlantic data interface signals are invalid and must be ignored. To determine whether new data has been received, the master must qualify the Val_o signal with the previous state of the Ena_i signal. The Rx buffers always provide all words of a packet on consecutive cycles (Val_o asserted) as long as Ena_i remains asserted during the packet. In addition, if Ena_i is asserted on the last word of a packet and the next packet is available, the next packet starts on the cycle immediately after the current packet completes.
Dav_o	Output	Data available. This signal functions as the Dav signal in the Atlantic interface specification with the HyperTransport MegaCore function as the slave source. If Dav_o is high, the buffer has at least one command/data packet available to be read. If this signal is not asserted, it indicates that there are no valid packets available to be read.
Ena_i	Input	Data transfer enable. This signal functions as the Ena signal in the Atlantic interface specification with the HyperTransport MegaCore function as a slave source. Ena_i is driven by the interface master and is used to control the flow of data across the interface. Ena_i behaves as a read enable from master to slave. When the slave observes Ena_i asserted on the RefClk rising edge, it drives, on the following RefClk rising edge, the Atlantic data interface signals and asserts Val_o. The master captures the data interface signals on the following RefClk rising edge.
BarHit_o[2:0]	Output	BAR match indication. For the Rx posted and Rx non-posted interfaces, this bus indicates which BAR the packet matched. 000 32-bit BAR0 or 64-bit BAR01 001 32-bit BAR1 010 32-bit BAR2 or 64-bit BAR23 011 32-bit BAR3 100 32-bit BAR4 or 64-bit BAR45 101 32-bit BAR5 110 Not used 111 Non-address packet This bus is valid only when Val_o and Sop_o are asserted. Because responses are not claimed due to address matches, this bus does not exist in the response buffer interface.

Rx Buffer Ordering

The circular Rx buffers allow the local-side application to read packets in the order in which they were received from the link. This implementation satisfies the HT specification ordering requirements within the same virtual channel and between the different channels.

If a response or non-posted packet with the `PassPW` bit reset is received, the HT specification requires that the packet not be processed until after the preceding posted packet is processed. The response and non-posted buffers are designed to enforce this requirement. For example, if a response packet with `PassPW` bit reset is received, the MegaCore function cannot read that packet from the response buffer until all preceeding posted packets are read from the Rx posted buffers. In this case, `RxRDav` is not asserted, and if the user interface asserts `RxREna`, `RxRVal` is deasserted to indicate invalid data until the preceeding posted packets are read from the posted buffer.

If `PassPW` is set on a received non-posted or response packet, the packets can be read out of the Rx buffers as soon as they are head of queue in their virtual channel, independent of any posted requests that may have arrived earlier.

☞ If a response packet with the `PassPW` bit set is received after another response packet with the `PassPW` bit reset, you cannot read the packet with the `PassPW` bit set until the previous one with the `PassPW` bit reset is read from the buffer. The same is true for non-posted packets.

Table 3–6 shows the Rx buffer ordering rules implemented by the HyperTransport MegaCore function.

Table 3–6. Rx Buffer Ordering

Row Pass Column		Posted Request		Non-Posted Request		Response	
Request Type	PassPW	Specification *(1)*	MegaCore Function	Specification *(1)*	MegaCore Function	Specification *(1)*	MegaCore Function
Posted	1	Yes/No	No	Yes	Yes *(2)*	Yes	Yes *(2)*
	0	No					
Non-posted	1	Yes/No	Yes *(2)*	Yes/No	No	Yes/No	Yes *(2)*
	0	No	No				
Response	1	Yes/No	Yes *(2)*	Yes	Yes *(2)*	Yes/No	No
	0	No	No				

Notes to Table 3–6:

(1) This column indicates the requirements as specified by the *HyperTransport I/O Link Specification Revision 1.03*.
No—The row request type may not pass the column request type.
Yes—The row request type must be allowed to pass the column request type to avoid deadlock conditions.
Yes/No—The row request type may pass the column request type but there is no requirement to do so.

(2) A "Yes" indicates that the MegaCore function allows the row to pass the column if the user application asserts the row `Ena_i` before the column `Ena_i`.

Rx Command/Data Buffer Timing Diagrams

Figure 3–10 shows a 32-byte (8 double-word) packet received across the Rx response buffer interface. In this example, the user-side logic asserts `Ena_i` in response to `Dav_o`. The Rx posted and Rx non-posted buffers behave in the same manner, except that posted and non-posted commands are typically 8 bytes. Thus, `Dat_o[63:32]` would be valid as well as `Sop_o`.

Figure 3–10. Single 32-Byte Read Response across Rx Interface, No Wait States Timing Diagram
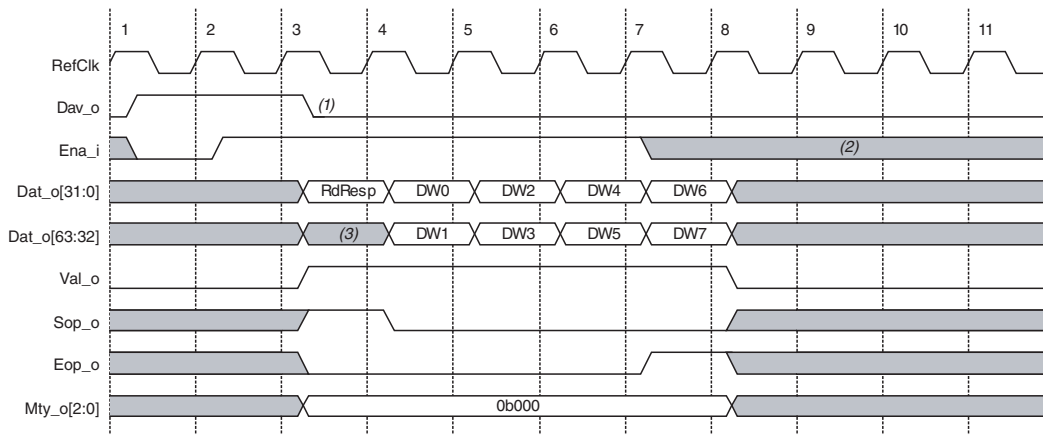
Notes to Figure 3–10:

(1) `Dav_o` is shown to deassert, indicating that the Rx buffer does not have any additional packets stored beyond the one that is currently being read.

(2) `Ena_i` is a Don't Care during this time because `Dav_o` is not asserted.

(3) `Dat_o[63:32]` is not valid during this time because the command is only a 4-byte read response.

Figure 3–11 is similar to Figure 3–10 except that `Ena_i` is always asserted so that the data transfer begins the clock cycle after `Dav_o` is asserted. Additionally, the example in Figure 3–11 is only 28 bytes long to show the behavior of `Mty[2:0]` during `Eop_o`.

Figure 3–11. 28-Byte Read Response across Rx Interface, No Wait States, Ena_i Asserted Timing Diagram
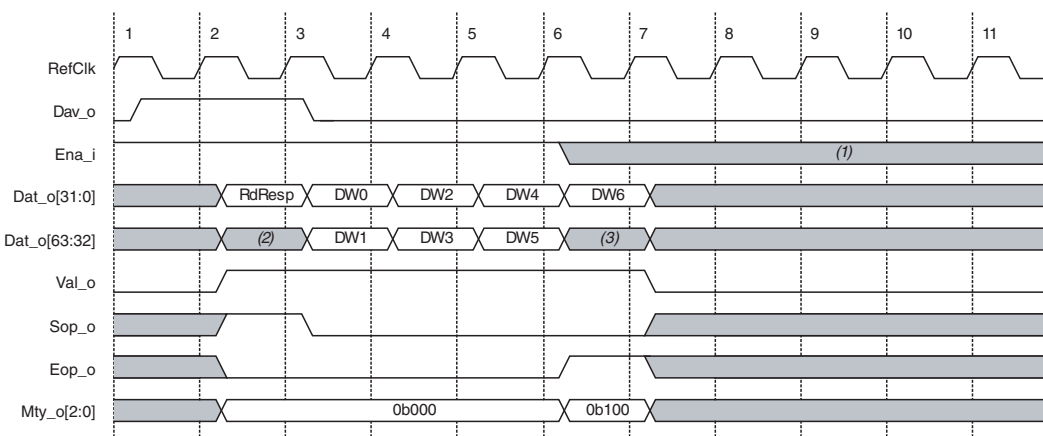
Notes to Figure 3–11:

(1) `Ena_i` is a Don't Care during this time because `Dav_o` is not asserted

(2) `Dat_o[63:32]` is not valid during this time because the command is only a 4-byte read response.

(3) `Dat_o[63:32]` is not valid during this time because the response is only seven double-words long.

Figure 3–12 shows a 32-byte write interrupted with user-inserted wait states the clock cycle after the command Sop_o = 1 is asserted.

Figure 3–12. 32-Byte Write with User-Inserted Wait States across Rx Interface Timing Diagram
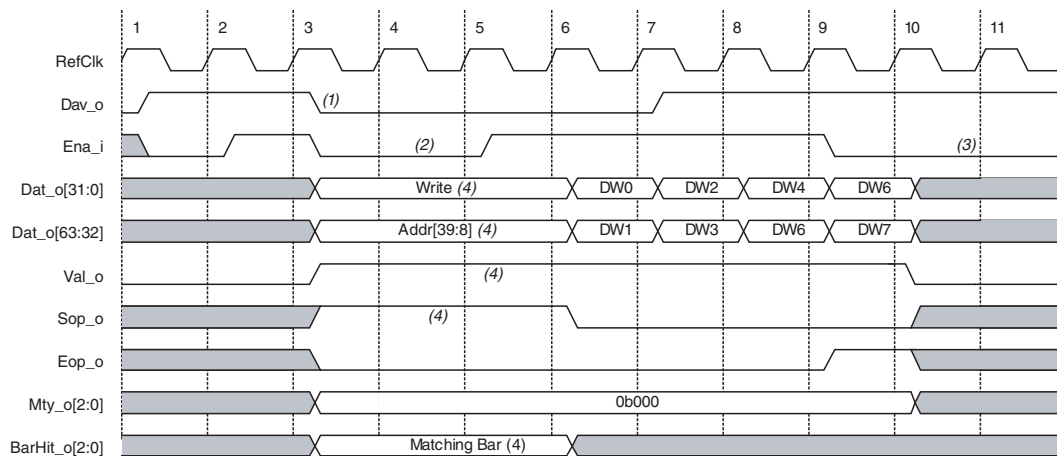
Notes to Figure 3–12:

(1) Dav_o is shown to deassert initially, indicating that the Rx buffer does not have any additional packets stored beyond the one that is currently being read.

(2) The user-side logic deasserts Ena_i to insert a wait state (for example to decode a command).

(3) Ena_i is low at this time so that the next command (indicated by Dav_o) is held off.

(4) Outputs are still valid, the user must qualify them with the previous Ena_i to determine if it is new data.

(5) BarHit_o is valid while Val_o and Sop_o are asserted.

Figure 3–13 shows a 38-byte read response, target done, and 16-byte read response transferred back to back across the Rx interface with no delays.

Figure 3–13. Streaming Responses Transferred across Rx Interface Timing Diagram
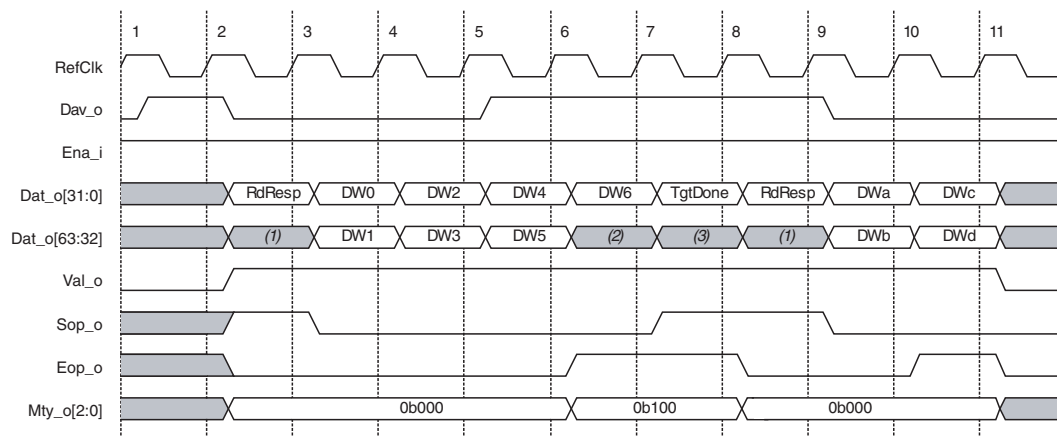
Notes to Figure 3–13:

(1) Dat_o[63:32] is not valid during this time because the command is only a 4-byte read response.

(2) Dat_o[63:32] is not valid during this time because the response is only seven double words long.

(3) Dat_o[63:32] is not valid during this time because the command is only a 4-byte target done.

Tx Command/Data Buffer Interfaces Operation

The Tx command/data buffers are temporary storage for all packets to be transmitted to the link. The Tx command/data buffer interfaces are slave sink Atlantic interfaces.

Each packet stored in the buffer contains a 32- or 64-bit command followed by its associated data, if there is any. The `Sop_i` signal indicates the start of packet, which must always be an HT command packet. Transmitted packets that do not have data are always one DWORD, and have `Eop_i` and `Sop_i` asserted.

HT command packets can be either 32 or 64 bits. 32-bit commands should be stored in the least significant bytes `Dat_i[31:0]`. The buffer automatically decodes the command, determines if it is 32 or 64 bits, and determines whether or not there is data associated with it. The buffer ignores `Mty` for the command word. It is illegal to have a non-zero value for the `Mty` signals at the command word if the command has data. The `Eop` signal indicates the end of the packet. When `Eop` is valid, the `Mty` signals are valid and indicate which bytes are invalid. Only the most significant four bytes can be invalid.

The `Mty` signals have only two valid values for the HyperTransport MegaCore function:

■ 100 indicates that `Dat_i[63:32]` is invalid

■ 000 indicates that all bytes are valid

After a packet is written to the buffer, the local-side application cannot retrieve that packet or prevent its transmission. A HT link warm or cold reset resets all buffer status bits and effectively destroys all data in the buffers. After the first word of a packet is written to the buffer, the entire packet must follow before a new packet is started. The transmitter circuitry begins transmission of a packet only after the entire packet is written into the buffer. This implementation allows the local application to use as much time as it needs to complete the packet.

Tx Command/Data Buffer Interface Signals

Table 3–7 describes the Tx buffer interface signals. Each of these signals has one of the following three types:

■ Tx response interface (HT MegaCore function signal names are prefixed with TxR)

■ Tx posted interface (HT MegaCore function signal names are prefixed with TxP)

■ Tx non-posted interface (HT MegaCore function signal names are prefixed with TxNp)

Because each channel uses similar interface signals, they are only described once. Any channels that have specific differences are noted in Table 3–7.

Table 3–7. Tx Command/Data Buffer Interface Signals   (Part 1 of 2)

Signal Name	Direction	Description
Dat_i[63:0]	Input	Data bus. This 64-bit input bus carries the data input to the command/data buffer. Data is loaded in little endian format (least significant byte is loaded first on Dat_i[7:0]).
Mty_i[2:0]	Input	Data byte empty. This bus indicates which bytes are invalid on Dat_i[63:0]. Because HT packets are DWORD aligned, there are only two valid encodings used by the HyperTransport MegaCore function. The following Mty_i signal values are acceptable: Mty_i = '000', all bytes on Dat_i[63:0] are valid. Mty_i = '100', Dat_i[63:32] are invalid. Invalid bytes can only be on the last word of a packet. It is illegal to have a non-zero value for the Mty_i signals for words other than the last word of the packet, even for 32-bit commands that have data, such as a read response. In this case, when Sop_i is asserted, Dat_i[63:32] is implicitly invalid. The first data bytes are placed on the Dat_i bus the cycle after the Sop_i. If the 32-bit command does not have data, the Mty_i signals should be set to '100' and Eop_i and Sop_i should be asserted.
Sop_i	Input	Start of packet. This signal must be high at the start of the packet. Start of packet is always loaded with least significant bytes loaded first. The first 64-bit word of a packet must always be the command. If the command is only 32 bits, it must be loaded on Dat[31:0] and the buffer ignores Dat_i[63:32].
Eop_i	Input	End of packet. This signal indicates the end of packet. Eop_i must be high at the end of the packet. When this signal is high, the Mty_i signals can indicate that Dat_i[63:32] are invalid.
WrRjct_o	Output	Write reject. This signal indicates that the local-side application's attempted write into the buffer has been rejected due to an error. The following errors cause the WrRjct_o signal to be asserted: ■ A write is attempted while Dav_o is low. ■ A write of the first word of a packet (command) was performed without Sop_i asserted. ■ A write is attempted that has more than 9 valid cycles. The maximum transfer is the HT command packet followed by a 64 byte HT data packet. If any of the above errors are detected, the write is rejected; the internal address counters and buffer status registers are reset to indicate that there were no writes after the last successful packet. Packets that were written into the buffer before the packet with the error are not affected and are transmitted to the link as usual. The WrRjct_o signal is intended for use as a debugging signal for simulation and early prototypes.

Table 3–7. Tx Command/Data Buffer Interface Signals   (Part 2 of 2)

Signal Name	Direction	Description
DatEna_i	Input	Data transfer enable. This signal functions as the `ena` signal in the Atlantic interface specification with the HyperTransport MegaCore function as a slave sink. `DatEna_i` is driven by the interface master and used to control the flow of data across the interface. `DatEna_i` behaves as a write enable.
Dav_o	Output	Data buffer is available. This signal functions as `dav` in the Atlantic interface. When `Dav_o` is high, it indicates that the buffer has sufficient storage for at least one complete packet including command and data.
		If there is room for at least two complete, full-size packets (9 64-bit words per packet) when a start of packet is written, `Dav_o` remains asserted during the packet so that another packet can be written immediately after the first. However, if there is only room for one full-size packet, `Dav_o` is deasserted after the start of the packet is written and reasserted again when the buffer has the room for at least one additional full-size packet.
		The local-side application must sample `Dav_o` asserted to write the start of the packet, but it should ignore `Dav_o` to write the subsequent words of the packet.
		When a single word packet (i.e. with `Sop_i` and `Eop_i` asserted in the same cycle) is written, it is typically not possible to start another packet in the following cycle. If there is initially only room for one additional packet, then `Dav_o` is deasserted in the cycle following the single word packet making it illegal to start writing a packet in that cycle. Typically the user application must write the single word packet in cycle one, sample `Dav_o` in cycle two and if `Dav_o` is still asserted, start writing a new packet in cycle 3.
		Special note for `TxRDav_o` only: Because the Tx response buffer also handles responses generated internally in the HyperTransport MegaCore function by the CSR and end-chain modules, the `TxRDav_o` signal can be deasserted at any time due to the internal usage. By default, the HyperTransport MegaCore function accepts response packets written by the user interface for one cycle after `TxRDav_o` is deasserted. If the local-side application requires multiple clock cycles between sampling `TxRDav_o` asserted and writing the start of packet (for example, due to a complex state machine or data pipeline), the window for accepting packets can be increased by changing the `TxRDavToSopDelay` configuration parameter.
		`TxPDav_o` and `TxNpDav_o` are only deasserted in response to user-written packets or reset.

Tx Buffer Ordering

Packets in the Tx buffers are transmitted in the same order they are written to the different buffers. For example, the response buffer transmits all packets in the order they are written into it. This implementation satisfies all ordering rule requirements in the HT specification except for response and non-posted packets that have the `PassPW` bit cleared. If a response or non-posted packet has the `PassPW` bit cleared, the HT specification requires that the response or non-posted packet must not be transmitted before the preceding posted packet.

The transmit packets have special handling for the case in which a response or non-posted packet has the `PassPW` bit reset. In this case, the transmit buffers enforce strict HT requirements. For example, if a response packet that has the `PassPW` bit reset is written to the transmit response buffer, the previously written posted packet is transmitted first. Because it is possible to write to the transmit posted buffer and transmit response buffer simultaneously, the previously written posted command is defined as the one that was started on or before the clock cycle preceding the start of packet of the response packet. The same is true in the case of non-posted packets.

Table 3–8 describes the Tx buffer ordering implemented by the HyperTransport MegaCore function.

Table 3–8. Tx Buffer Ordering

Row Pass Column		Posted Request		Non-Posted Request		Response	
Packet	PassPW	Specificaton *(1)*	MegaCore Function	Specification *(1)*	MegaCore Function	Specification *(1)*	MegaCore Function
Posted Request	1	Yes/No	No	Yes	Yes *(2)*	Yes	Yes *(2)*
	0	No					
Non-Posted Request	1	Yes/No	Yes *(2)*	Yes/No	No	Yes/No	Yes *(2)*
	0	No	No				
Response	1	Yes/No	Yes *(2)*	Yes	Yes *(2)*	Yes/No	No
	0	No	No				

Notes to Table 3–8:

(1) This column indicates the requirements as specified by the *HyperTransport I/O Link Specification Revision 1.03*.
No—The row request type may not pass the column request type.
Yes—The row request type must be allowed to pass the column request type to avoid deadlock conditions.
Yes/No—The row request type may pass the column request type but there is no requrement to do so.

(2) A "Yes" indicates that the row may pass the column subject to available buffer credits and internal scheduling.

Tx Command/Data Buffer Interface Timing Diagrams

Figure 3–14 shows a 32-byte read response being transferred to the HyperTransport MegaCore function across the Tx interface.

Figure 3–14. Single 32-Byte Read Response across Tx Interface, No Wait States Timing Diagram
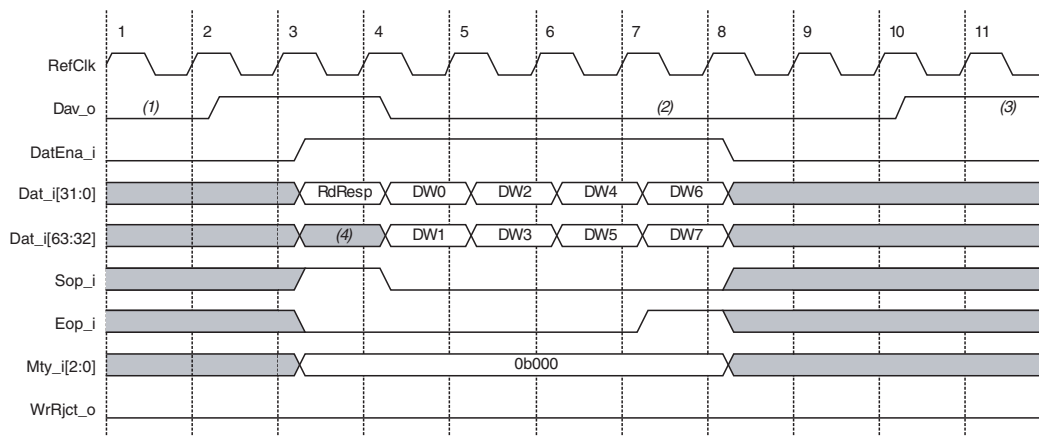
Notes to Figure 3–14:

(1) In this example, `Dav_o` starts out low. The clock cycle after `Dav_o` is asserted, the user application asserts `DatEna_i`. In other cases, `Dav_o` may be high for many clock cycles and the user application does not assert `DatEna_i` until it has data available.

(2) `Dav_o` goes low in this example to indicate the TxResp buffer is full and cannot accept a subsequent command. However, the user application still asserts `DatEna_i` to complete the current command.

(3) Some time later `Dav_o` is reasserted, indicating another command can be accepted.

(4) `Dat_o[63:32]` is not valid in this example because the command is only a read response.

Figure 3–15 shows a 28-byte write being transferred to the HyperTransport MegaCore function across the Tx Interface. The user logic inserts a wait state during the transfer.

Figure 3–15. User Wait States on the Tx Interface Timing Diagram
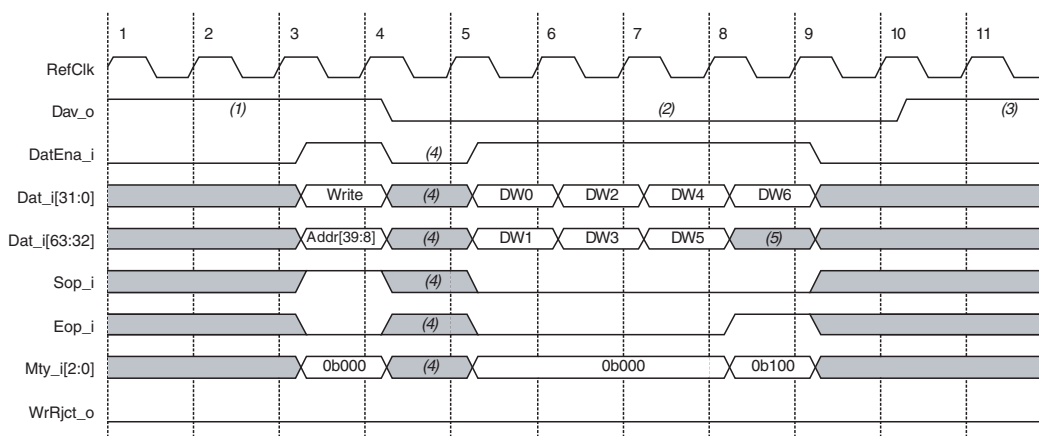
Notes to Figure 3–15:

(1) In this example, `Dav_o` is high for many clock cycles and the user application asserts `DatEna_i` when it has data available.

(2) `Dav_o` goes low in this example to indicate the TxResp buffer is full and cannot accept a subsequent command. However, the user application still asserts `DatEna_i` to complete the current command.

(3) Some time later `Dav_o` is reasserted, indicating another command can be accepted.

(4) The user inserts a wait state in the transfer by deasserting `DatEna_i`; the other inputs are Don't Cares in this clock cycle.

(5) `Dat_i[63:32]` is not valid because this example is only seven double-words.

Figure 3–16 shows two commands transferred back-to-back across the Tx response buffer interface. A 36-byte read response is transferred followed by a 16-byte read response.

Figure 3–16. Two Streamed Commands on Tx Response Buffer Interface Timing Diagram
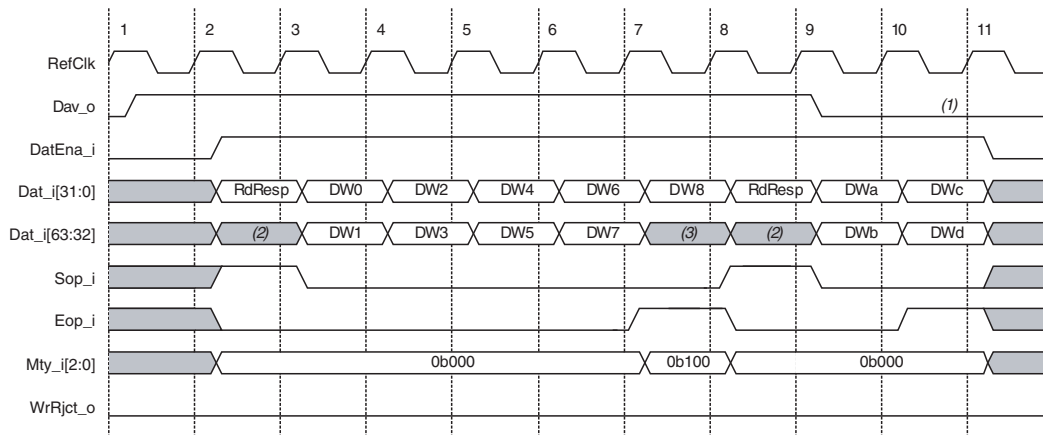
Notes to Figure 3–16:

(1) Dav_o goes low in this example to indicate the TxResp buffer is full and cannot accept a subsequent command. However, the user application still asserts DatEna_i to complete the current command.

(2) Dat_i[63:32] is not valid because the read response command is only 4 bytes. Mty_i[2:0] is still 3'b000 because Eop_i is not asserted.

(3) Dat_i[63:32] is not valid because the the data is an odd number of DWORDS in length. Mty_i[2:0] is 3'b100 because Eop_i is asserted.

Figure 3–17 shows the how the Tx Buffer WrRjct_o output can be triggered by attempting to write to a Tx buffer when Dav_o is not asserted.

Figure 3–17. Tx Buffer Write Reject Timing Diagram
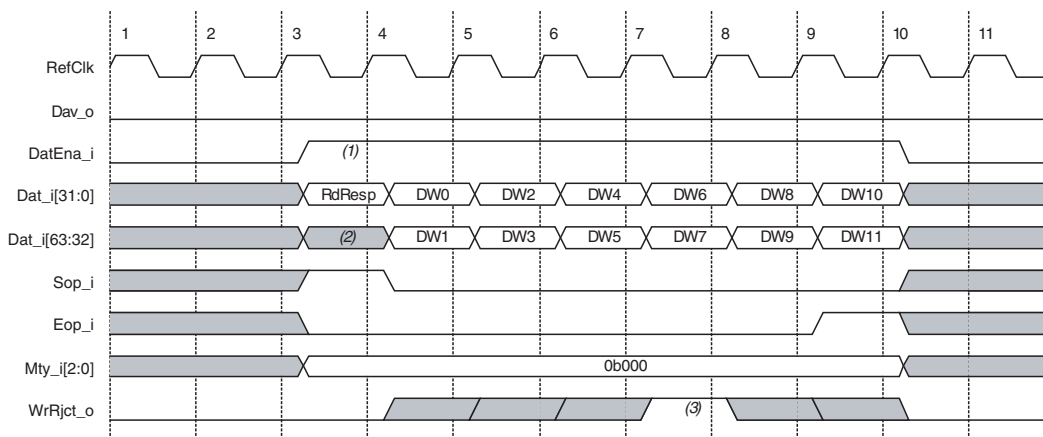
Notes to Figure 3–17:

(1) In this example, DatEna_i is asserted without Dav_o being asserted, leading to an overrun of the Tx buffer.

(2) Dat_o[63:32] is not valid in this example because the command is only a read response.

(3) WrRjct_o is asserted because DatEna_i was asserted with Dav_o. The exact clock cycle in which WrRjct_o is asserted varies due to internal buffer conditions.

Internally, the HyperTransport MegaCore writes to the Tx response Buffer to write CSR access responses as well as end-chain responses. To do so, it deasserts `TxRDav_o` so that the user-side application does not write to the buffer at the same time. By default, it allows one clock cycle after the deassertion of `TxRDav_o` to accept further writes from the user side, however, the time that user writes are allowed can be increased by setting the `TxRDavToSopDelay` configuration parameter to a higher value. This implementation only applies to the Tx response buffer. The Tx posted and non-posted buffers deassert `Dav_o` only in response to the user writing a packet.

Figure 3–18 shows the deassertion of `TxRDav_o` and a subsequent user attempt to write to the Tx response buffer with the `TxRDavToSopDelay` parameter set to the default value of one.

Figure 3–18. HyperTransport MegaCore Function Internal Deassertion of TxRDav_o Timing Diagram
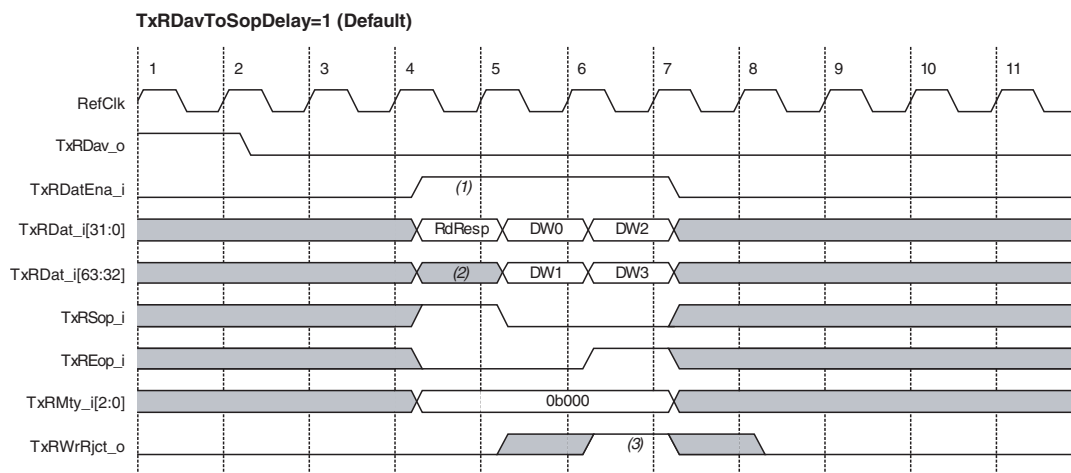
Notes to Figure 3–18:

(1) In this example, `TxRDatEna_i` is asserted two clock cycles after `Dav_o` is deasserted due to an internal use of the Tx response buffer (`TxRDatEna_i` was not asserted). This Tx buffer write is deasserted due to a conflict with the internal usage.

(2) `TxRDat_o[63:32]` is not valid in this example because the command is only a read response.

(3) `TxRWrRjct_o` is asserted because `TxRDatEna_i` was asserted with `TxRDav_o`. The exact cycle in which `TxRWrRjct_o` is asserted varies due to internal buffer conditions.

Figure 3–19 shows the effect of writing to the Tx response buffer with the `TxRDavToSopDelay` parameter set to five.

Figure 3–19. Tx Response Buffer with the TxRDavToSopDelay Set to 5 Timing Diagram
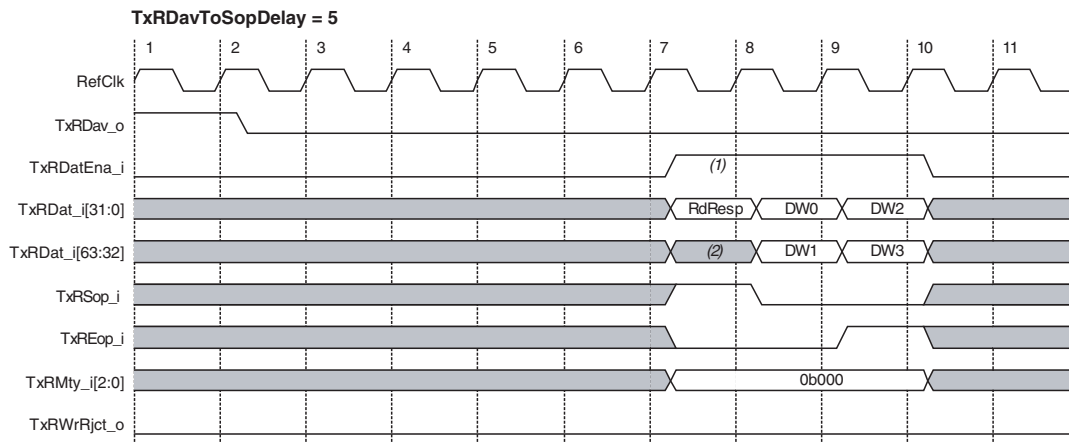
Notes to Figure 3–19:

(1) In this example, `TxRDatEna_i` is asserted five clock cycles after `Dav_o` is deasserted due to an internal use of the Tx response buffer (`TxRDatEna_i` was not asserted). This Tx buffer write is not rejected because the `TxRDavToSopDelay` parameter is set to five.

(2) `TxRDat_o[63:32]` is not valid during this time because the command is only a read response.

User Interface System Signals

Table 3–9 lists the user interface system signals.

Table 3–9. User Interface System Signals

Signal	Direction	Description
RxLnkClkD4	Output	Clock generated by the Rx PLL. This clock is created based on the transmit clock of the HT device connected to the MegaCore function's receiver, and is four times slower than its source. When using the **Shared Rx/Tx/Ref Clock** option, the user interface logic must use this clock.
RxLnkClkD4Locked_o	Output	`RxLnkClkD4` lock indicator. This signal is asserted when the Rx PLL is locked.
RefClkWrmRst	Output	Warm reset signal synchronized to the reference clock input (`RefClk`).
RefClkCldRst	Output	Cold reset signal synchronized to the reference clock input (`RefClk`).
RefClk	Input	Reference clock. This user clock source drives the MegaCore function's protocol module. When using the **Shared Rx/Tx/Ref Clock** option, this input is not available.
RxPllAreset_i	Input	Resets the HyperTransport Rx PLL used in all clocking modes.
TxPllAreset_i	Input	Resets HyperTransport Tx PLL used in the shared Ref/Tx clocking mode. This input, which can be tied low, is not used in the other modes.

Using PLL Resets

If the input clocks of the fast PLLs used in the HyperTransport MegaCore function stop and restart, the output phase relationships of the PLL output clocks may not meet their designed values unless the PLL is reset using these inputs after the clock has restarted.

During normal operation of the HyperTransport link, clocks may change from the initial 200 MHz frequency to a faster frequency. If the `RxClk_i` input to the function stops during this change, the Rx PLL must be reset. The HyperTransport interface does not provide sufficient indication to directly generate the PLL reset for this situation. However, a Verilog HDL reference design that generates the HyperTransport Rx PLL reset at the appropriate time is installed with the HyperTransport MegaCore function.

The Verilog HDL reference design code that generates the HT Rx PLL reset at the appropriate time can be found in your Altera MegaCore IP library installation at *<path>*/**ht/example/pll_reset/ht_with_pll_reset.v**

The `RxPllAreset_i` and `TxPllAreset_i` inputs can be tied low to provide the PLL behavior that was present in HyperTransport MegaCore function v1.3.0 and earlier.

CSR Interface Signals

All required status information in the CSR module is provided to the local side through the signals shown in Table 3–10.

Table 3–10. CSR Interface Signals

Signal Name	Direction	Description
`CsrCmdReg[15:0]`	Output	Command register.
`CsrStatReg[15:0]`	Output	Status register.
`CsrCapReg[15:0]`	Output	Capability command register.
`CsrCapLnk0CtrlReg[15:0]`	Output	Capability link 0 control register.
`CsrCapLnk1CtrlReg[15:0]`	Output	Capability link 1 control register.
`CsrCapLnk0CfgReg[15:0]`	Output	Capability link 0 configuration register.
`CsrCapLnk1CfgReg[15:0]`	Output	Capability link 1 configuration register.
`CsrCapFtrReg[7:0]`	Output	Capability feature register.
`CsrCapLnk0ErrReg[3:0]`	Output	Capability link 0 error register.
`CsrCapLnk1ErrReg[3:0]`	Output	Capability link 1 error register.
`CsrCapErrHndlngReg[15:0]`	Output	Capability error handling register.
`Bar0Reg[31:0]`	Output	Base address register 0.
`Bar1Reg[31:0]`	Output	Base address register 1.
`Bar2Reg[31:0]`	Output	Base address register 2.
`Bar3Reg[31:0]`	Output	Base address register 3.
`Bar4Reg[31:0]`	Output	Base address register 4.
`Bar5Reg[31:0]`	Output	Base address register 5.
`RespErr`	Input	The user application logic asserts this signal to indicate that an unexpected response was received. Asserting this signal causes the appropriate CSR bit to be set.
`SignaledTabrt`	Input	The user application logic asserts this signal to indicate that a target abort was signaled (the target issued a response with a non-NXA error). Asserting this signal causes the appropriate CSR bit to be set.

Buffer Overflow Indicator Signals

Table 3–11 describes the buffer overflow indicator signals. When an a Rx buffer overflows due to a failure in the HyperTransport link flow control mechanism, the overflow error bit is set in the CSR HT capability link error 0 register. These signals indicate which virtual channel buffer had the overflow. These signals are most useful in simulation and typically do not need to be connected in the user design.

Table 3–11. Buffer Overflow Indicator Signals

Signal Name	Direction	Description
ClmdRCmdBufOvrFlw_Err_o	Output	Rx Response Command/Data buffer overflow.
ClmdPCmdBufOvrFlw_Err_o	Output	Rx Posted Command/Data buffer overflow.
ClmdNpCmdBufOvrFlw_Err_o	Output	Rx Non-Posted Command/Data buffer overflow.

CSR Module

The CSR module contains all of the configuration space registers and handles write requests, read requests and response accesses to the CSR space. Table 3–12 and Table 3–13 show the full CSR register map. Table 3–13 shows the primary interface capabilities block format.

Table 3–12. CSR Register Map Without Primary Interface Capabilities Block

Hyper Transport Technology Device Header				
Address	Byte 3	Byte 2	Byte1	Byte 0
00	Device ID *(3)*		Vendor ID *(3)*	
04	Status *(4)*		Command *(4)*	
08	Class Code *(3)*			Revision ID *(3)*
0C	BIST *(1)*	Header Type *(3)*	Latency Timer *(1)*	Cache Line *(1)*
10	BAR0 *(4)*			
14	BAR1 *(4)*			
18	BAR2 *(4)*			
1C	BAR3 *(4)*			
20	BAR4 *(4)*			
24	BAR5 *(4)*			
28	CardBus CIS Pointer *(1)*			
2C	Subsystem ID *(3)*		Subsystem Vendor ID *(3)*	
30	Expansion ROM Base Address *(2)*			
34	Reserved *(1)*			Capabilities Pointer *(3)*
38	Reserved *(1)*			
3C	Max Latency *(1)*	Min Grant *(1)*	Interrupt Pin *(1)*	Interrupt Line *(2)*
+18	Reserved *(1)*		Mem Limit Upper *(2)*	Mem Base Upper *(2)*

Notes to Table 3–12:

(1) The configuration register is not used by the HT specification. The register is a read-only register, and Read returns 0s.

(2) The register is not supported in the HyperTransport MegaCore function. The register is a read-only register and Read returns 0s.

(3) The register is a read-only register in the HyperTransport MegaCore function.

(4) The register is read/write in the HyperTransport MegaCore function.

Table 3–13 shows the format of the primary interface capabilities block. The block is located at the address specified in the Capabilities Pointer field in the CSR map shown in Table 3–12.

Table 3–13. Primary Interface Capabilities Block Format

Address	HyperTransport Slave/Primary Interface Capabilities Block Format			
	Byte 3	**Byte 2**	**Byte1**	**Byte 0**
+00	Command (1)		Capabilities Pointer (2)	Capabilities ID (2)
+04	Link Config 0 (1)		Link Control 0 (1)	
+08	Link Config 1 (1)		Link Control 1 (1)	
+0C	Link Frequency Cap 0 (2)		Link Error 0 (1) / Link Freq 0 (1)	Revision ID (2)
+10	Link Frequency Cap 1 (2)		Link Error 1 (1) / Link Freq 1 (1)	Feature (1)
+14	Error Handling (1)		Enumeration Scratch Pad (4)	
+18	Reserved (3)		Mem Limit Upper (4)	Mem Base Upper (4)

Notes to Table 3–13:

(1) The register is read/write in the HyperTransport MegaCore function.

(2) The register is a read-only register in the HyperTransport MegaCore function.

(3) The configuration register is not used by the HT specification. The register is a read-only register, and Read returns 0s.

(4) The register is not supported in the HyperTransport MegaCore function. The register is a read-only register and Read returns 0s.

The following sections explain the entries in Table 3–12 and Table 3–13.

HT Device Header

The following sections describe the HT device header.

Vendor ID, Device ID (Offset 0x00)

Vendor ID and device ID are read-only, as defined in the PCI specification. They are parameters in the HyperTransport MegaCore function.

PCI Command Register (Offset 0x04)

Table 3–14 describes the PCI command register.

Table 3–14. PCI Command Register (Offset 0x04)   (Part 1 of 2)

Bits	Name	Register Type	Reset Used	Description
0	I/O Space Enable	R/W	Warm	Must be set for the device to claim non-compatibility accesses to I/O address space
1	Memory Space Enable	R/W	Warm	Must be set for the device to claim non-compatibility accesses to memory address space.
2	Bus Master Enable	R/W	Warm	Must be set for the device to issue memory or I/O requests onto the HT chain.
7:3	Reserved			

Table 3–14. PCI Command Register (Offset 0x04)   (Part 2 of 2)

Bits	Name	Register Type	Reset Used	Description
8	SERR# Enable	R/W	Warm	If set, the device floods all outgoing links with sync packets.
15:9	Reserved			

PCI Status Register (Offset 0x06)

Table 3–15 describes the PCI status register.

Table 3–15. PCI Status Register (Offset 0x06)

Bits	Name	Register Type	Reset Used	Description
3:0	Reserved			
4	Capabilities List	R/O	N/A	Hardwired to 1 to indicate that there is a capabilities list present.
10:5	Reserved			
11	Signaled Target Abort	R/C	Cold	Set by device when a response with non-NXA error occurs (generated response with non-NXA error).
12	Received Target Abort	R/C	Cold	Indicates that the device received a target abort (received response with non-NXA error).
13	Received Master Abort	R/C	Cold	Indicates that the device has received a master abort (received response with NXA set).
14	Signaled System Error	R/C	Cold	Set by the device to indicate that it has signaled SERR (generated sync flood).
15	Reserved			

Revision ID, Class Code (Offset 0x08)

The revision ID and class code are read-only as defined in the PCI specification. They are parameters in the HyperTransport MegaCore function.

Header Type (Offset 0x0E)

The header type is read-only as defined in the PCI specification. The default is 0.

Base Address Registers (Offsets 0x10, 0x14, 0x18, 0x1C, 0x20, 0x24)

The HT BAR definition is similar to the PCI specification's BAR definition. The incoming address is first compared to the I/O address map to determine if it is in the memory range or the I/O range. For memory ranges, the BAR can be either 64 or 32 bits. If the BAR is 64 bits, the incoming address must be zero-extended to 64-bits before being compared to the BAR value. All 6 BARs are supported with up to 3 64-bit BARs or up to 6 32-bit BARs. The only valid 64-bit BAR pairs are:

- BAR0 and BAR1
- BAR2 and BAR3
- BAR4 and BAR5

Subsystem Vendor ID, Subsystem ID (Offset 0x2C)

The subsystem vendor ID and subsystem ID are read-only as defined in the PCI specification. They are parameters in the HyperTransport MegaCore function.

Capabilities Pointer (Actual Offset 0x34)

This register is read-only register for the HyperTransport MegaCore function. It points to 0x40 in the HyperTransport MegaCore function.

HT Capabilities Block

The HyperTransport MegaCore function uses a slave/primary interface capabilities block format as defined in the HT specification. Table 3–13 shows this block format.

Capabilities ID (Actual Offset 0x40)

This register is read only as defined by the HT specification and is hardwired to 0x08.

Capabilities Pointer (Actual Offset 0x41)

This register is read only as defined by the HT specification. In the HyperTransport MegaCore function, it is hardwired to 0x00, indicating no additional capability blocks.

HT Capability Command Register (Actual Offset 0x42)

Table 3–16 describes the HT capability command register.

Table 3–16. HT Capability Command Register Offset: Capability Pointer + 0x02 (Actual Offset: 0x42)

Bits	Name	Register Type	Reset Used	Description
4:0	Base Unit ID	R/W	Warm	Indicates the lowest number UnitID belonging to the device.
9:5	Unit Count	R/O	N/A	Number of UnitID occupied by this device. The MegaCore function does not occupy more than one UnitID.
10	Master Host	R/O *(1)*	N/A	Because a device can have up to two links (link 0, link 1) this bit indicates which one links to the host bridge. The HyperTransport MegaCore function hardwires this bit to 0 to indicate that link 0 is always connected to the host bridge. Additionally, the CSR module can only be accessed via link 0.
11	Default Direction	R/O *(2)*	N/A	Indicates which directions to which the request should be sent. This bit is hardwired to 0 to indicate that all requests should be sent to the host bridge.
12	Drop on Uninit	R/W	Cold	Drop on uninitialized link. See *(Note 3)* for details.
15:13	Capabilities Type	R/O	N/A	Indicates the type of capabilities. This register is hardwired to 000 to indicate that the capabilities type is slave/primary.

Notes: to Table 3–16:

(1) This register is defined as a read/write in the HT specification. The HyperTransport MegaCore function does not allow both links in a tunnel to write to the CSR. Therefore, this bit is always hardwired to 0 to indicate that link 0 is the one connected to the host bridge.

(2) This bit is defined as Read/Write in the HT specification so that the direction is programmable. The HyperTransport MegaCore function only supports one direction.

(3) This bit determines what happens to packets loaded in the transmitter while init complete and end chain bits are cleared. Refer to the *HyperTransport I/O Link Specification Revision 1.03* for a description of this functionality.

Link Control Registers (Actual Offsets 0x44, 0x48)

The HyperTransport MegaCore function cannot be software configured. The user must assume that link zero is connected to the host bridge. Additionally, link 1 is connected to the device end chain. The bits of link control 0 register are described in the following table. The link control one register is hardwired to 0x0050, indicating that the device is an end-chain device. Table 3–17 describes the HT capability link control registers.

Table 3–17. HT Capability Link Control Register Bits Offset: Capability Pointer + 0x04 (Actual Offset: 0x44)

Bits	Name	Register Type	Reset Used	Description
0	Reserved			
1	CRC Flood Enable	R/W	Warm	The link generates a sync flood error if this bit is set. Otherwise, the CRC checker logs the errors.
2	CRC Start Test	HW to 0	N/A	The HyperTransport MegaCore function does not support the CRC test sequence. This bit is hardwired to zero.
3	CRC Force Error	R/O	N/A	This bit forces bad CRC generation. The HyperTransport MegaCore function does not support this mode.
4	Link Failure	R/C	Cold	Set this bit when a failure is detected and causes a sync flood. These errors include CRC, protocol error, or overflow error. This bit is hardwired to one for end-chain applications.
5	Initialization Complete	R/O	Warm	Set to 1 automatically when the low-level initialization is complete.
6	End of Chain	R/O	Warm	Indicates that a link is not part of the logical HT chain. This bit is parameterized. Link zero always has its bit hardwired to zero; link one is hardwired to one.
7	Transmitter Off	R/O	N/A	Allows shutting down the transmitter. When this bit is set, outputs do not toggle. This bit is hardwired to zero.
11:8	CRC Error	R/C	Cold	One of these bits is set when a CRC error is detected. Errors are detected and reported on per byte basis. Because the 8-bit HyperTransport MegaCore function only implements one byte, bits 11,10, and 9 are hardwired to 0.
12	Isoc Enable	R/O	N/A	Controls whether the isochronous flow control mode is supported or not. Because the HyperTransport MegaCore function does not support the isochronous flow control mode, this bit is hardwired to zero.
13	LDTSTOP# Tristate Enable	R/O	N/A	Controls whether the link tristates the signals during a disconnect state of LDTSTOP# sequence. Because the HyperTransport MegaCore function does not support power management, this bit is hardwired to 0.
14	Extended CTL Time	R/O	N/A	If this bit is set during the initialization sequence, CTL is asserted for 50 μs after the point where both Tx and Rx assert CTL. The MegaCore function only supports the 16-bit time after both Tx and Rx assert CTL, therefore, this bit is hardwired to zero.
15	Reserved			

Link Configuration Registers (Actual Offsets 0x46, 0x4A)

There are two link configuration registers in every device, one for each link. The unused register for the unused link is reserved. The HyperTransport MegaCore function always uses link configuration register 0 in the case of an end-chain instantiation. Table 3–18 describes the HT capability link configuration registers.

Table 3–18. HT Capability Link Configuration Register Bits Offset: Capability Pointer + 0x06 (Actual Offset: 0x46)

Bits	Name	Register Type	Default	Reset Used	Description
2:0	Max Link Width In	R/O	0	N/A	Indicates the maximum width of the incoming side of the HyperTransport MegaCore function. The value is hardwired to 000 to indicate an 8-bit maximum.
3	Double Word Flow Control In	R/O	0	N/A	Indicates that the device is capable of double word-based data buffer flow control. This bit is hardwired to zero.
6:4	Max Link Width Out	R/O	0	N/A	Indicates the maximum width of the Tx port of the HyperTransport MegaCore function. This bit is hardwired to indicate 8 bits.
7	Double Word Flow Control Out	R/O	0	N/A	Indicates that Tx is capable of double word flow control. This bit is hardwired to zero.
10:8	Link Width In	R/O	0	N/A	Indicates the actual width used by the link. This bit is hardwired to 000 because the MegaCore function does not have an automatic sizing module.
11	Double Word Flow Control In Enable	R/O	0	N/A	This bit indicates whether the double word flow control mode will be enabled. This bit is hardwired to zero because the MegaCore function does not provide this mode.
14:12	Link Width Out	R/O	0	N/A	Indicates the actual width used by the link. This bit is hardwired to 000 because the MegaCore function does not have automatic sizing module.
15	Double word Flow control Out Enable	R/O	0	N/A	This bit indicates whether the double word flow control mode will be enabled. This bit is hardwired to zero because the MegaCore function does not provide this mode.

Revision ID (Actual Offset 0x4C)

This register shows the revision of the specification used to design the link interface. This register is hardwired with 0x13.

Link Frequency Registers (Actual Offsets 0x4D, 0x51)

There are two registers of this type, one for each link. In the HT specification this register is intended to control the frequency of the link transmit clock. However in the HyperTransport MegaCore function, the link transmit frequency is always four times the `RxLnkClkD4` frequency (or when using the Shared Ref/Tx Clock option it is four times `RefClk`) and the value of this register has no effect inside the MegaCore function.

☞ This register is a read/write register for compatibility with configuration software.

Link Error Registers (Actual Offsets 0x4D, 0x51)

There are two registers of this type, one for each link. In the HyperTransport MegaCore function, register 0 defines the link connected to the host bridge. For end-chain applications, only register zero is used. Table 3–19 describes the HT capability link error registers.

Table 3–19. HT Capability Link Error 0 Register Offset: Capability Pointer + 0x0D (Actual Offset: 0x4D)

Bits	Name	Register Type	Default	Reset Used	Description
3:0	Link Frequency Register	R/W	0	Cold	Intended to control the frequency of the link transmit clock, but has no effect in the HyperTransport MegaCore function.
4	Protocol Error	R/C	0	Cold	Indicates that a protocol error was detected. The HyperTransport MegaCore function reports a protocol error when the RxCTL_i signal changes on other than a 4-byte boundary.
5	Over Flow Error	R/C	0	Cold	Indicates that a packet was received but there was no room in the Rx packet buffer. Indicates a failure in the flow control mechanism.
6	End of Chain Error	R/C	0	Cold	Indicates a posted request of response has been given to this transmitter when it is an end chain.
7	CTL Timeout	R/W	0	Warm	Indicates how long CTL may be low before a device indicates a protocol error. 0 indicates one millisecond; 1 indicates 1 second.

Link Frequency Capability Registers (Actual Offsets 0x4E, 0x52)

Each link frequency capability register is a 16-bit read only register that indicates the frequency capability of the link. Two registers are defined, one for each link. For end-chain applications, only register 0 is used. Table 3–20 and Table 3–21 describe the link frequency registers.

Table 3–20. HT Capability Link Frequency 0 Capability Register Bits Offset Capability Pointer + 0x0E (Actual Offset: 0x4E)

Bits	Name	Register Type	Default	Reset Used	Description
15:0	Supported Frequencies	R/O	See Table 3–21.	N/A	Indicates the supported frequencies of operation. The default value is set based on the HT_RX_CLK_PERIOD Parameter.

Table 3–21. Link Frequency 0 Capability Register Default Value

HT_RX_CLK_PERIOD	Register Value	Supported Frequencies (MHz)
5000	0x0001	200
3333	0x0003	200, 300
2500	0x0007	200, 300, 400
2000	0x000F	200, 300, 400, 500

Feature Capability Register (Actual Offset 0x50)

This register indicates which optional features the device supprts. Table 3–22 describes the HT capability feature capability register.

Table 3–22. HT Capability Feature Capability Register Bits Offset: CapabilityPointer + 0x10 (Actual Offset: 0x50)

Bits	Name	Register Type	Default	Reset Used	Description
0	Isochronous Mode	R/O	0	N/A	Indicates that the HyperTransport MegaCore function is not capable of isochronous mode.
1	LDTSTOP#	R/O	0	N/A	Indicates that the HyperTransport MegaCore function does not support LDTSTOP# protocol.
2	CRC Test Mode	R/O	0	N/A	Indicates that the HyperTransport MegaCore function does not support CRC test mode.
3	Extended CTL Time Required	R/O	0	N/A	Indicates that the HyperTransport MegaCore function does not require extended CTL time.
6:4	Reserved				
7	Extended Register Set	R/O	1	N/A	Indicates that the HyperTransport MegaCore function implements the extended error handling register. It does not implement the enumeration scratch pad register or the memory base/limit upper registers.

Enumeration Scratch Pad Register (Actual Offset 0x54)

This register is not implemented in the HyperTransport MegaCore function. It is hardwired to zero.

Error Handling Register (Actual Offset 0x56)

Table 3–23 shows the HT capability error handling register bits.

Table 3–23. HT Capability Error Handling Register Bits Offset: Capability Pointer + 0x16 (Actual Offset: 0x56)   (Part 1 of 2)

Bits	Name	Register Type	Default	Reset Used	Description
0	Protocol Error Flood Enable	R/W	0	Warm	When set, this bit results in a sync flood when a protocol error is detected
1	Overflow Error Flood	R/W	0	Warm	When set, this bit results in a sync flood when an overflow error is detected.
2	Protocol Error Fatal Enable	R/O	0	N/A	When set, this bit results in a fatal error interrupt whenever the protocol error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
3	Overflow Error Fatal Error	R/O	0	N/A	When set, this bit results in a fatal error interrupt whenever the overflow error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
4	End of Chain Error Fatal Enable	R/O	0	N/A	When set, this bit results in a fatal error interrupt whenever the end-chain error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.

Table 3–23. HT Capability Error Handling Register Bits Offset: Capability Pointer + 0x16 (Actual Offset: 0x56)   (Part 2 of 2)

Bits	Name	Register Type	Default	Reset Used	Description
5	Response Error Fatal Enable	R/O	0	N/A	When set, this bit results in a fatal error interrupt whenever the response error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
6	CRC Error Fatal Enable	R/O	0	N/A	When set, this bit results in fatal error interrupt whenever one of the CRC error bits is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
7	System Error Fatal Enable	R/O	0	N/A	This bit is for host interfaces only.
8	Chain Fail	R/O	0	Warm	This bit is set if a sync flood is detected or a sync-flood generating error.
9	Response Error	R/C	0	Cold	This bit indicates that the link has received a response error.
10	Protocol Error NonFatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever the protocol error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
11	Overflow Error Non-Fatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever overflow error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
12	End of Chain Error Non-Fatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever end-of-chain error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
13	Response Error Non-Fatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever the response error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
14	CRC Error Non-Fatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever any of the CRC error bits are asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.
15	System Error Non-Fatal Enable	R/O	0	N/A	When set, this bit results in a nonfatal error interrupt whenever the response error bit is asserted. This bit is hardwired to zero because the HyperTransport MegaCore function does not support error interrupts.

Memory Base/Limit Upper Registers (Actual Offset 0x58)

These registers are defined for bridge functions and are hardwired to zero.

OpenCore Plus Time-Out Behavior

OpenCore Plus hardware evaluation can support the following two modes of operation:

■ *Untethered*—the design runs for a limited time.

■ *Tethered*—requires a connection between your board and the host computer. If tethered mode is supported by all megafunctions in a design, the device can operate for a longer time or indefinitely.

All megafunctions in a device time out simultaneously when the most restrictive evaluation time is reached. If there is more than one megafunction in a design, a specific megafunction's time-out behavior may be masked by the time-out behavior of the other megafunctions.

☞ For MegaCore functions, the untethered timeout is one hour; the tethered timeout value is indefinite.

Your HyperTransport MegaCore function variation is forced into a cold reset state when the hardware evaluation time expires.

🐾 For more information about OpenCore Plus hardware evaluation, refer to "OpenCore Plus Evaluation" on page 1–3 and *AN 320: OpenCore Plus Evaluation of Megafunctions*.

Introduction

You configure the HyperTransport MegaCore function entirely using the IP Toolbench wizard interface. This section describes the parameters that are modified by settings you make in the wizard interface.

For more information about the IP Toolbench wizard, see "Step 1: Parameterize" on page 2–5.

Parameter Lists

Table A–1 through Table A–4 list the parameters in each tab of the wizard. The parameter names listed here are identical to those used in the HDL variation files created by the IP Toolbench.

Device Family and Read Only Registers

Table A–1 lists the parameters modified by settings on the **Device Family & Read-Only Registers** tab of the IP Toolbench wizard.

Table A–1. Device Family & Read-Only Registers Parameters

Parameter	Default	Available Setting(s)	Description
INTENDED_DEVICE_FAMILY	Stratix	Stratix II GX, Stratix II, Stratix, Stratix GX	Specifies the family of the target device. If this does not match the device selected for the Quartus II project, the Quartus II software reports an error during compilation.
DeviceID	0x0000	Any 16-bit value	The value to be returned when the CSR device ID register is read.
VendorID	0x1172	Any 16-bit value	The value to be returned when the CSR vendor ID register is read. Should be set to the users PCI-SIG defined vendor ID. (The default value is Altera's vendor ID.)
ClassCode	0x00_0000	Any 24-bit value	The value to be returned when the CSR class code register is read. Should be set based on the PCI-SIG defined class encoding for the type of device being built.
RevisionID	0x00	Any 8-bit value	The value to be returned when the CSR revision ID register is read.
SubSystemID	0x0000	Any 16-bit value	The value to be returned when the CSR subsystem ID register is read.
SubSystemVendID	0x1172	Any 16-bit value	The value to be returned when the CSR subsystem vendor ID register is read. (The default value is Altera's vendor ID.)

Base Address Registers

Table A–2 lists the parameters modified by settings on the **Base Address Registers** tab of the IP Toolbench wizard.

Table A–2. BAR Parameters

Parameter	Default	Available Setting(s)	Description
Bar0AndMask	0xFFF0_0000	Any 32-bit value	The AndMask defines the writable bits of the specified BAR. For 32 bit memory space BARs, only bits 31-10 can be writable, i.e., the minimum decodable memory space is 1KBytes. Only BAR1, BAR3, and BAR5 can be the upper BARs of a 64 bit BAR. In this case, all bits can be writable.
Bar1AndMask	0x0	Any 32-bit value	
Bar2AndMask	0x0	Any 32-bit value	
Bar3AndMask	0x0	Any 32-bit value	
Bar4AndMask	0x0	Any 32-bit value	
Bar5AndMask	0x0	Any 32-bit value	
Bar0MatchMask	0xFFF0_0000	Any 32-bit value	The MatchMask specifies which bits of the specified BAR should be used for matching addresses. In all cases, the MatchMask should be identical to the AndMask.
Bar1MatchMask	0x0	Any 32-bit value	
Bar2MatchMask	0x0	Any 32 -bit value	
Bar3MatchMask	0x0	Any 32-bit value	
Bar4MatchMask	0x0	Any 32-bit value	
Bar5MatchMask	0x0	Any 32-bit value	
Bar0OrMask	0x8	Any 4-bit value	The OrMask specifies which low-order type bits of the BAR should be set to 1. ■ Bit[0] – I/O space ■ Bit[1] – Reserved ■ Bit[2] – 64 bit BAR (only valid on BAR0, BAR2, and BAR4) Only valid when Bit [0]=0 ■ Bit[3] – Prefetchable On the upper BAR of a 64 bit BAR, all of these bits must be 0.
Bar1OrMask	0x0	Any 4-bit value	
Bar2OrMask	0x0	Any 4-bit value	
Bar3OrMask	0x0	Any 4-bit value	
Bar4OrMask	0x0	Any 4-bit value	
Bar5OrMask	0x0	Any 4-bit value	
CLAIM_COMPAT	0x0	0 or 1	Compatibility Request Claiming occurs based on whether the CLAIM_COMPAT parameter is set to zero or 1. When the value of this parameter is zero, the HyperTransport MegaCore function does not claim Read and Write requests that have the Compat bit set to 1, posted requests are ignored and nonposted requests have a NXA error response returned. When the CLAIM_COMPAT parameter is set to the value 1, the requests are claimed and presented on the Rx user interfaces of the function. Setting the value to 1 can be used to implement an x86 Subtractive Decode device. The CLAIM_COMPAT parameter is set by the **Claim All Read And Write Requests With Compat Bit Equal To** '1' box on the **Base Address Registers** tab.

Clocking Options

Table A–3 lists the parameters modified by settings on the **Clocking Options** tab of the IP Toolbench wizard.

Table A–3. Clocking Options Parameters

Parameter	Default	Available Setting(s)	Description
CLKDMN	3	1, 2, 3	This parameter selects the clocking option for the HyperTransport MegaCore function. Values have the following meanings: 1 – **Shared Rx/Tx/Ref Clock** option 2 – **Shared Ref/Tx Clock** option 3 – **Shared Rx/Tx Clock** option Refer to "Clocking Options" on page 3–7 for a discussion of these options.
HT_RX_CLK_PERIOD	2500	2000, 2500, 3333, 5000	HT Rx link clock period in picoseconds. This parameter sets the Rx PLL parameters for the correct Rx frequency. Set this parameter to the highest operating frequency. The PLL still locks at the slower frequencies.
HT_TX_CLK_PERIOD	2500	2000, 2500, 3333, 5000	HT Tx link clock period in picoseconds. This parameter sets the Tx PLL parameters for the correct Tx frequency. Set this parameter to the highest operating frequency. The PLL still locks at the slower frequencies. This parameter is only used in the **Shared Ref/Tx Clock** option. In the other clock options, the Tx PLL is shared with the Rx PLL.

Advanced Settings

Table A–4 lists parameters modified by settings on the **Advanced Settings** tab of the IP Toolbench wizard.

Table A–4. Advanced Settings Parameters   (Part 1 of 2)

Parameter	Default	Available Setting(s)	Description
RxPBufNum	8	8, 16	Selects either 8 or 16 Rx posted buffers. Selecting 16 Rx posted buffers may increase throughput but also consumes additional resources.
RxNpBufNum	4	4, 8	Selects either 4 or 8 Rx non-posted buffers. Selecting 8 Rx non-posted buffers may increase performance but also consumes additional resources.

Table A–4. Advanced Settings Parameters   (Part 2 of 2)

Parameter	Default	Available Setting(s)	Description
RxRBufNum	4	4, 8	Selects either 4 or 8 Rx response buffers. Selecting 8 Rx response buffers may increase performance but also consumes additional resources.
TxRDavToSopDelay	1	1 - 15	This parameter defines the number of cycles after the MegaCore function de-asserts TxRDav_o during which it accepts an additional local-side write to the Tx response buffer. This parameter is useful if the local-side application takes more than one cycle from sampling TxRDav_o asserted until it first starts writing the packet to the buffer. However, increasing this value does increase the time it takes for CSR or end-chain handler responses to be transmitted, as well as any local-side response that immediately follows.

Introduction

This section provides general guidelines on how to assign the HyperTransport pins for Stratix devices. Many rules differ based on the device and package you use, and not all variations are documented here.

⚠️ **CAUTION**
You must compile your pin assignments in the Quartus II software version 9.0 or later to verify they are correct before you commit the pin assignments to the board layout. You should make sure the pin assignments will also work for any device you may migrate to later, or any clock option you may consider using in the future.

The I/O pins that support the HyperTransport differential I/O standard are on the sides of Stratix devices, specifically, I/O banks B1, B2, B5, and B6.

Guidelines

The general guidelines for HyperTransport I/O pin assignments are described below:

- If your design uses a 400-MHz HyperTransport link clock, `RxClk_i` should use one of these PLL clock input pin pairs that support the highest input clock rates:

    - `CLK0p/CLK0n` - Fast PLL1 - I/O Bank B2

    - `CLK2p/CLK2n` - Fast PLL2 - I/O Bank B1

    - `CLK9p/CLK9n` - Fast PLL3 - I/O Bank B6

    - `CLK11p/CLK11n` - Fast PLL4 - I/O Bank B5

If the `CLK1`, `CLK3`, `CLK8`, and `CLK10` inputs are used and drive the same respective PLLs, the HyperTransport link frequency is limited to 300 MHz.

👣 See the *DC & Switching Characteristics* chapter of the *Stratix Device Handbook* for the most up to date information about the maximum input clock rates.

- All of the I/O pins for the receive interface (`RxCAD_i[7:0]`, `RxCTL_i`, and `RxClk_i`) must be in the same I/O bank, so that the high-speed de-serialization circuitry can be connected to the same fast PLL.

- All of the I/O pins for the transmit interface (`TxCAD_o[7:0]`, `TxCTL_o`, and `TxClk_o`) must be in the same I/O bank, so that the high-speed serialization circuitry can be connected to the same fast PLL.

- If you are using the **Shared Rx/Tx/Ref Clock** option or the **Shared Rx/Tx Clock** option, all of the transmit and receive interface pins must be in the same I/O bank, because they all share a single fast PLL.

■ If you are using the **Shared Ref/Tx Clock** option, which requires separate transmit and receive PLLs, and you are using a device with only four PLLs (for example, the EP1S10, EP1S20, EP1S25, EP1S30 devices in 780-pin BGA packages, or the EP1S40 in 780-pin BGA packages), the Rx and Tx interfaces must be in separate I/O banks because these devices have only one PLL per differential I/O bank.

■ If you are using the **Shared Ref/Tx Clock** option, which requires separate transmit and receive PLLs, and you are using a Stratix device that has 8 PLLs (e.g., the EP1S30 and EP1S40 devices in most packages, EP1S60, and EP1S80), the transmit and receive interfaces can be placed in the same I/O bank because there are two PLLs per differential I/O bank. You can also create an I/O pin out that supports any clocking option with these devices. To support any clock option, the chosen CAD and CTL I/O pins must be clustered in the center of the I/O bank so that they can be driven at a high speed from either the center or corner PLL. It is essential to test your pin out with Quartus II compilation for all clock options you might use before finalizing the pin out for board layout.

■ If your HyperTransport interface is running at 300 MHz (600 Mbps) or 400 MHz (800 Mbps), use only the pins listed as supporting a DIFFIO speed of "HIGH" in the Stratix device pin information. This information is available on the Altera website as part of the *Stratix Device Handbook*.

General Description

Figure C–1 shows the high-level diagram of a example design that shows simple usage of all of the HyperTransport MegaCore function interfaces. This design is an SDRAM memory controller with a DMA engine that transfers data between the HyperTransport interface and the SDRAM. In this example, the DMA engine is programmed by a host processor on the HyperTransport link. The DMA engine performs the requested data movement and the host processor reads the status back from the DMA engine register.

Figure C–1. Example Design
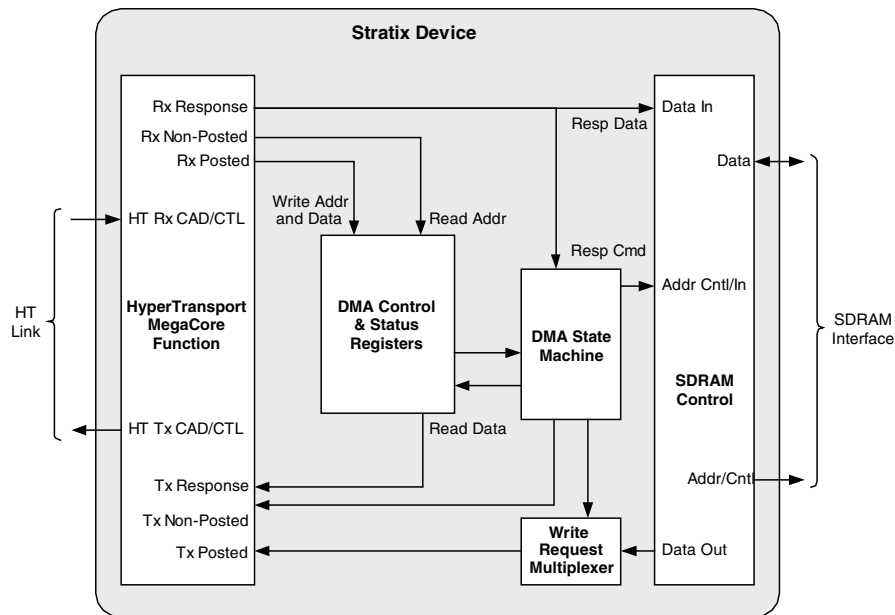
For a data movement that reads from the SDRAM and writes to host CPU memory through HyperTransport, the following sequence of events occurs:

1. The host CPU programs the DMA control registers using HyperTransport posted writes. The writes come from the Rx posted interface and update the DMA control registers.

2. The DMA state machine issues an SDRAM read request to the SDRAM control. When the SDRAM read data returns, the DMA state machine generates the HyperTransport post write request command, which is then multiplexed ahead of the data and written in the HyperTransport Tx posted buffer.

3. Step 2 repeats until the requested DMA transfer completes and all of the SDRAM read data is written to the HyperTransport Tx posted buffer.

4. The CPU reads the DMA status registers using an Rx non-posted read request and subsequent HyperTransport Tx read response. The response informs it that the operation has completed and it can program another DMA operation. HyperTransport ordering rules ensure that the status register read response does not pass the final DMA posted write, and the host CPU can be confident that all of the DMA data has been written to host memory.

For a data movement that reads data from host CPU memory and writes to the SDRAM the following sequence of events occurs:

1. The host CPU programs the DMA control registers using HyperTransport posted writes. The writes come from the Rx posted interface and update the DMA control registers.

2. The DMA state machine issues a read request to the HyperTransport interface through the Tx non-posted buffer. To hide the latency of accessing memory on the HyperTransport host CPU complex, the DMA state machine streams many read requests without waiting for the responses. Each read request uses a unique source tag in the HyperTransport request packet. (The source tag is set by the DMA state machine, not by the HyperTransport MegaCore function.) A single read request may not straddle a 64-byte boundary in the HyperTransport address space.

3. When the HyperTransport read responses begin arriving, the DMA state machine reads the read response command packet from the Rx response buffer and uses the returned source tag to determine the local SDRAM address. The SDRAM address and write control is transferred to the SDRAM control. The read response data packet is read from the Rx Response buffer and passed to the SDRAM control. SDRAM control performs the SDRAM write.

4. Steps 2 and 3 repeat until the requested DMA transfer completes and all of the read response data is returned and written to SDRAM.

5. The CPU reads the DMA status registers using an Rx non-posted read request and subsequent HyperTransport Tx read response. The read response informs it that the operation has completed and it can program another DMA operation.

Revision History

The following table shows the revision history for this user guide.

Date and Document Version	Changes Made	Summary of Changes
March 2009 v9.0	Maintenance release	Maintenance release
November 2008 v8.1	Maintenance release	Maintenance release
May 2008 v8.0	Updated the MegaCore IP version, and incorporated the following modifications: ■ Updated Table 3–12 to clarify behavior of some read-only registers ■ Described Quartus II IP File (**.qip**) ■ Fixed Features to clarify that this MegaCore function supports low-swing differential signaling with 100-ohm differential impedance	Maintenance release and minor text updates
October 2007 v7.2	Maintenance release	—
May 2007 v7.1	Maintenance release; updated release information	—
December 2006 v7.0	Maintenance release	—
December 2006 6.1	In Chapter 2, updated the MegaCore IP version, and incorporated the following documentation errata: ■ Setting Virtual Pin Attributes Recommended ■ Enabling Clock Latency Recommended for All Variations	—
October 2005 v1.4.0	Added Stratix® II GX support	—
February 2004 v1.3.0 rev 1	Added support for Altera® Stratix II devices, IP functional simulation models, and the OpenCore Plus evaluation feature	—

How to Contact Altera

For the most up-to-date information about Altera products, refer to the following table.

Contact (1)	Contact Method	Address
Technical support	Website	www.altera.com/support
Technical training	Website	www.altera.com/training
	Email	custrain@altera.com
Product literature	Website	www.altera.com/literature

Contact *(1)*	Contact Method	Address
Non-technical support (General)	Email	nacomp@altera.com
(Software Licensing)	Email	authorization@altera.com

Note to table:

(1)   You can also contact your local Altera sales office or sales representative.

Typographic Conventions

This document uses the following typographic conventions.

Visual Cue	Meaning
Bold Type with Initial Capital Letters	Command names, dialog box titles, checkbox options, and dialog box options are shown in bold, initial capital letters. Example: **Save As** dialog box.
bold type	External timing parameters, directory names, project names, disk drive names, file names, file name extensions, and software utility names are shown in bold type. Examples: **\qdesigns** directory, **d:** drive, **chiptrip.gdf** file.
Italic Type with Initial Capital Letters	Document titles are shown in italic type with initial capital letters. Example: *AN 75: High-Speed Board Design.*
Italic type	Internal timing parameters and variables are shown in italic type. Examples: t_{PIA}, $n + 1$. Variable names are enclosed in angle brackets (< >) and shown in italic type. Example: *<file name>*, *<project name>*.**pof** file.
Initial Capital Letters	Keyboard keys and menu names are shown with initial capital letters. Examples: Delete key, the Options menu.
"Subheading Title"	References to sections within a document and titles of on-line help topics are shown in quotation marks. Example: "Typographic Conventions."
`Courier type`	Signal and port names are shown in lowercase Courier type. Examples: `data1`, `tdi`, `input`. Active-low signals are denoted by suffix `n`, e.g., `resetn`. Anything that must be typed exactly as it appears is shown in Courier type. For example: `c:\qdesigns\tutorial\chiptrip.gdf`. Also, sections of an actual file, such as a Report File, references to parts of files (e.g., the AHDL keyword `SUBDESIGN`), as well as logic function names (e.g., `TRI`) are shown in Courier.
1., 2., 3., and a., b., c., etc.	Numbered steps are used in a list of items when the sequence of the items is important, such as the steps listed in a procedure.
■ ■	Bullets are used in a list of items when the sequence of the items is not important.
✓	The checkmark indicates a procedure that consists of one step only.
☞	The hand points to information that requires special attention.
⚠ CAUTION	A caution calls attention to a condition or possible situation that can damage or destroy the product or the user's work.
⚠ WARNING	A warning calls attention to a condition or possible situation that can cause injury to the user.
↵	The angled arrow indicates you should press the Enter key.
👣	The feet direct you to more information about a particular topic.